


Image Digitization: 3 Stages


1) Transforming the **3D** world into **2D** image

- Perspective Projection (Optics, Continuous)



2) Sampling the Image Plane

- Finite number of **Pixels**




3) Quantizing the color/gray-level


- Finite number of colors (e.g. 8 bits per color)



Pinhole Camera (Camera Obscura (Latin) = Dark Room)




- No Image is generated when we place an image plane in the world



- To create an image, each image location should get a ray from a single point. This is done by blocking all rays except one...

Pinhole Camera (Camera Obscura (Latin) = Dark Room)

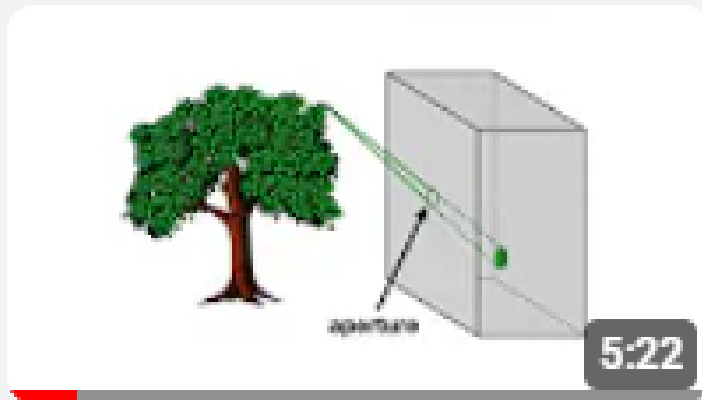


- Pinhole model:
 - Captures pencil of rays – all rays through a single point
 - The point (pinhole) is called Center of Projection (COP)
 - The image is formed on the Image Plane
 - Focal length f is distance from COP to Image Plane

5 Minute Video Clips

- 12) Perspective projection
- 13) Perspective projection: Part 2 – the math!


12



Perspective projection in 5 minutes

Graphics in 5 Minutes • 24K views • 2 years ago

13




Perspective projection in 5 minutes: Part 2 -- the math!


Graphics in 5 Minutes • 8.3K views • 2 years ago

Perspective Projection

- Transforming the 3D world (X, Y, Z) into 2D image (x, y)
 - Continuous Perspective Projection (optics)
 - All rays pass through one point (f = focal length)




Simple case: Aligned
World axis (X, Y, Z)
and Image axis (x, y)



$$\frac{Y}{Z} = \frac{y}{f}$$

Similar
Triangles

$$x = \frac{f}{Z} X$$
$$y = \frac{f}{Z} Y$$




Perspective Projection

- World to Camera Transformation


$$x = \frac{f}{Z} X \qquad y = \frac{f}{Z} Y$$

- Only when world axis (X, Y, Z) and Image axis (x, y) are aligned:
 - X is parallel to x
 - Y is parallel to y
 - Same units
- In the general case, there is a transformation matrix between world axis and camera axis.



Summary: First Stages of Image Acquisition

Analog Light



Spatial Sampling to Pixels

- Sampling the Image Plane
 - Finite number of **Pixels**
 - Do we always want maximum number of pixels?




256 lines



64 lines



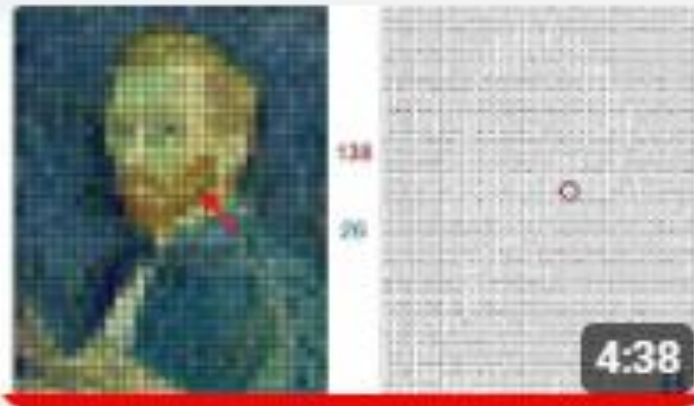
16 lines



8 lines

5 Minute Video Clips

- 4) Images in 5 minutes





Images in 5 minutes: The Case of the Splotched Van Gogh, Part 1

Graphics in 5 Minutes • 4.4K views • 2 years ago


RGB Inside the Camera

Anatomy of the Active Pixel Sensor Photodiode



Bayer Filter


- In 1975, Bruce Bayer invents the color filter array, used in most digital camera.
- $\frac{1}{4}$ pixels detect red; $\frac{1}{4}$ pixels blue; $\frac{1}{2}$ pixels green;
- The camera invents 2 missing colors in pixels. How?



- Demosaicing: Invents missing colors
- Many methods, mostly proprietary
- A possible (bad) method:
 - Average 2 or 4 neighbors

Color/Gray-level Quantization

- Quantizing the color/gray-level
 - Finite number of colors



256 Levels



8 Levels



4 Levels



2 Levels

Digital Pictures



- A Matrix of numbers (Greylevel image)
- A Matrix of triplets (RGB Color, etc.)

2	4	5	6	7	8	9	10	11	12	11	10	9	8	7
3	5	6	7	8	9	10	11	12	13	12	11	10	9	8
4	6	7	8	9	10	11	12	13	14	13	12	11	10	9
5	7	8	9	10	11	12	13	14	15	14	13	12	11	10
6	8	9	10	11	12	13	14	15	16	15	14	13	12	11
7	9	10	11	12	13	14	15	16	17	16	15	14	13	12
8	10	11	12	13	14	15	16	17	18	17	16	15	14	13
9	11	12	13	14	15	16	17	18	19	18	17	16	15	14
10	12	13	14	15	16	17	18	19	20	19	18	17	16	15
9	11	12	13	14	15	16	17	18	19	18	17	16	15	14
8	10	11	12	13	14	15	16	17	18	17	16	15	14	13
7	9	10	11	12	13	14	15	16	17	16	15	14	13	12
6	8	9	10	11	12	13	14	15	16	15	14	13	12	11
5	7	8	9	10	11	12	13	14	15	14	13	12	11	10
4	6	7	8	9	10	11	12	13	14	13	12	11	10	9
3	5	6	7	8	9	10	11	12	13	12	11	10	9	8
2	4	5	6	7	8	9	10	11	12	11	10	9	8	7
1	3	4	5	6	7	8	9	10	11	10	9	8	7	6

Color Spaces

RGB (Camera, Projector - Add), CMYK (Print -Subtract), YIQ (TV)

Projector:
Add
colors



Ink:
Absorb
colors



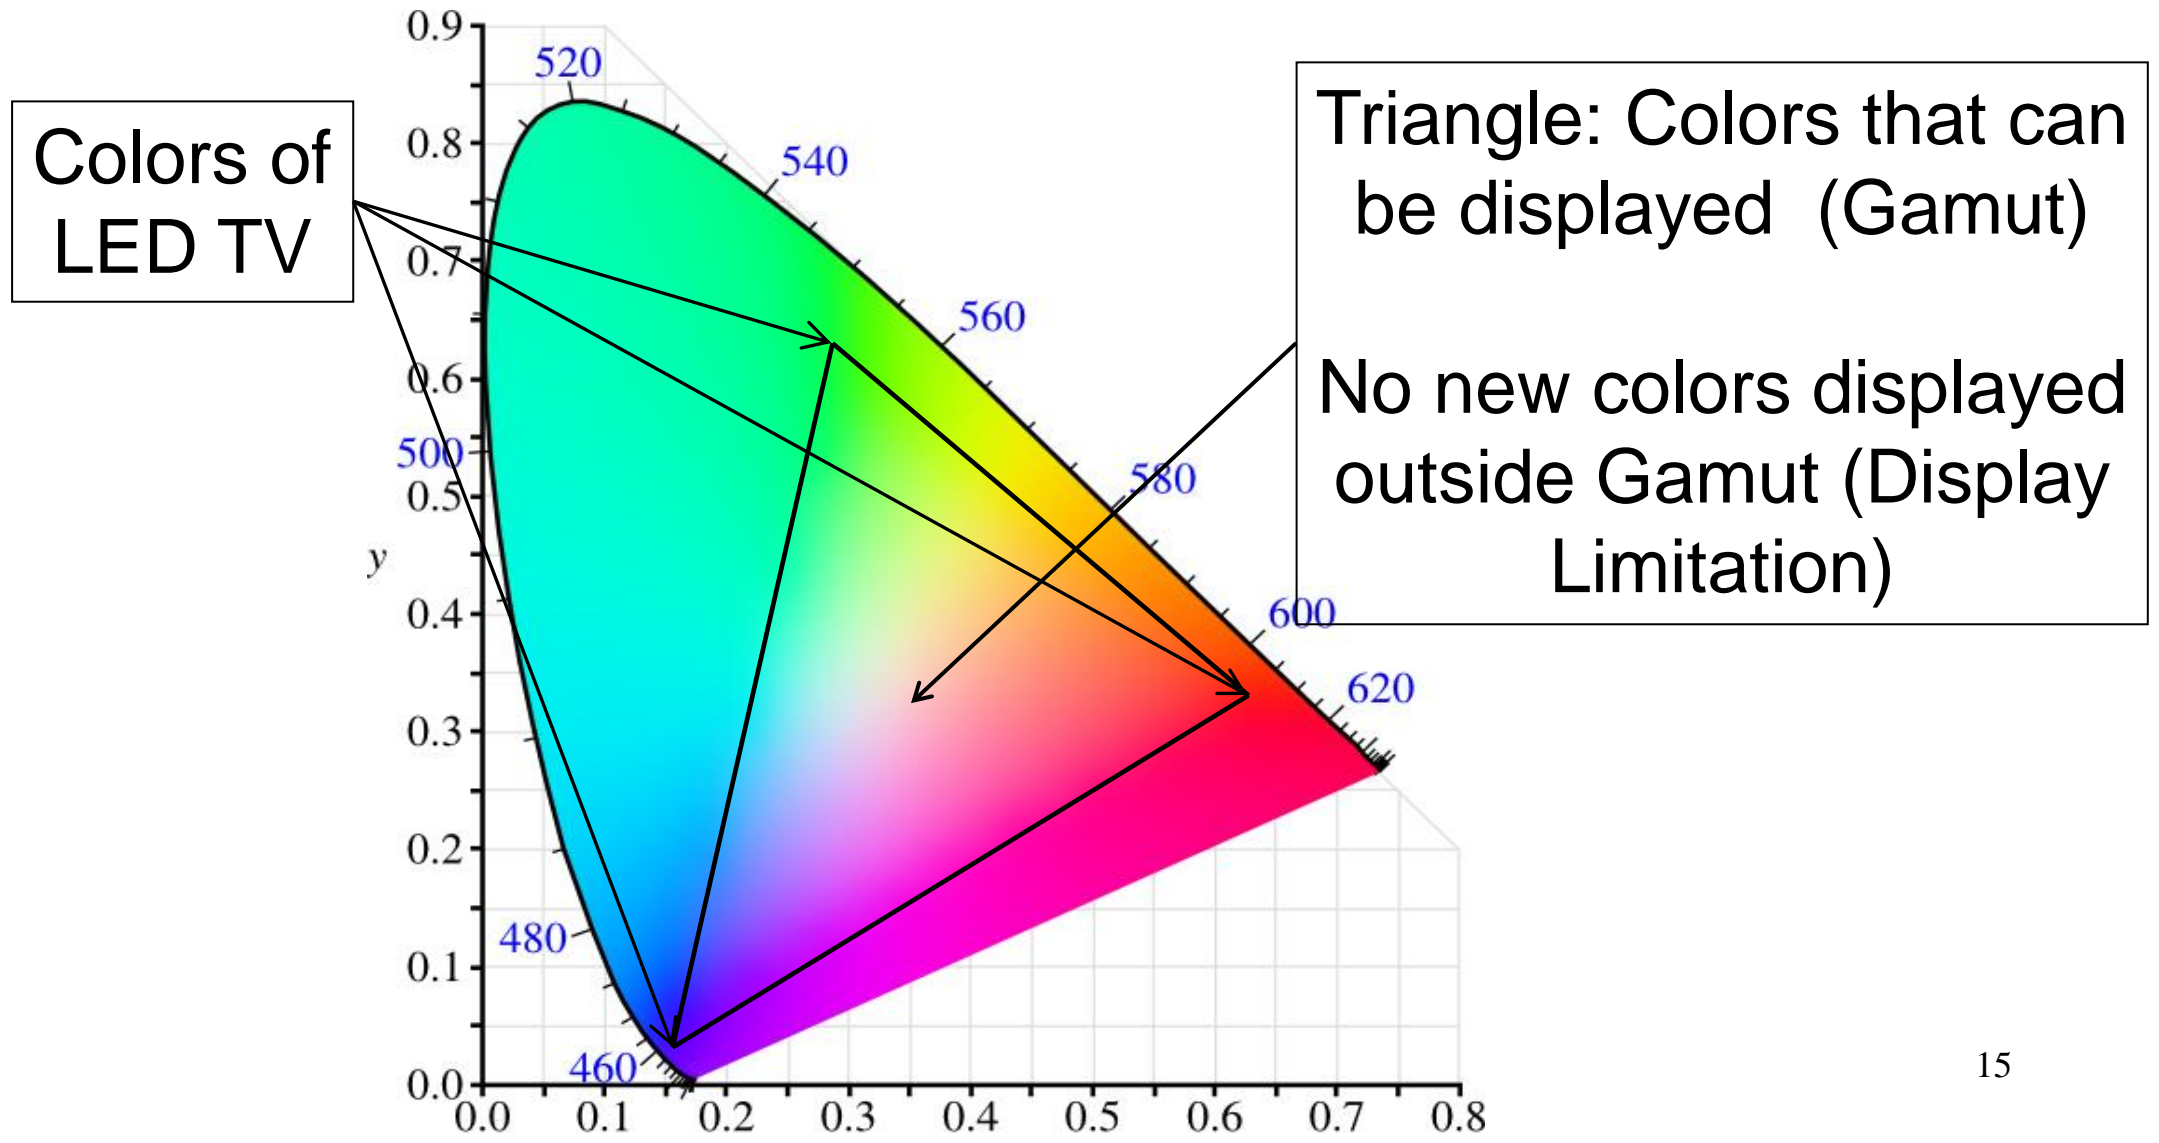
For Color to B/W TV

Y - Luminance


$$\begin{bmatrix} Y \\ I \\ Q \end{bmatrix} = \begin{bmatrix} 0.299 & 0.587 & 0.114 \\ 0.596 & -0.275 & -0.321 \\ 0.212 & -0.523 & 0.311 \end{bmatrix} \begin{bmatrix} R \\ G \\ B \end{bmatrix}$$

CIE Chromaticity Diagram (1931)

Boundary: Spectral Colors (Single Wavelength)



Point Operations



- New pixel value $g(x,y)$ based on the input pixel value $f(x,y)$
- E.g. $g(x,y)=255-f(x,y)$ (Negative)
- Operation depends only on pixel value (No location...)

Point Operations

f




g



- In general, $g(x, y) = T(f(x, y))$


- $T(u) = 255 - u$

Negative



- $T(u) = \begin{cases} 0 & \text{if } u < 127 \\ 1 & \text{if } u > 127 \end{cases}$

Threshold




Point Operations - γ Correction

Gamma Correction is used to overcome non-linear responses of camera, display, and eyes

$$T(u) = Max \cdot \left(\frac{u}{Max}\right)^\gamma$$

When $Max=1$: $T(u) = u^\gamma$




$\gamma = 1$



$\gamma = 1/2.2$

Eye Sensitivity


- Eye sensitivity as a function of frequency



Point Transformation - Look Up Table (LUT)

- LUT efficiently Represent transformations L from N to N (or to other). E.g.

- $L(0)=0$
- $L(1)=0$
- $L(2)=0$
- $L(3)=2$
- $L(4)=4$
- $L(5)=6$
- $L(6)=7$
- $L(7)=7$



Point Operation with LUT

- Stretch




In	0	1	2	3	4	5	6	7
Out	0	0	0	2	4	6	7	7

- Threshold

In	0	1	2	3	4	5	6	7
Out	0	0	0	0	7	7	7	7


- Negative

In	0	1	2	3	4	5	6	7
Out	7	6	5	4	3	2	1	0




The Histogram

- Frequency counting of gray levels or colors
- Analogous to PDF – Probability Density Function



Pixel Count	1	2	4	6	8	9	10	8	7	6	5	8	16	6	4	0
Grey Level	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16

Common Histogram has 256 Grey Levels



0

255

Count: 1920000

Min: 0

Mean: 118.848




Max: 251

StdDev: 59.179


Mode: 184 (30513)

Histogram is Invariant to Pixel Location


- Different pictures can give same histograms
- All three pictures below have same histogram. What is it?



Histogram & Exposure




OK





Over Exposed

JPG Compression Effects





- Original Image




- Image after JPG compression


Color Histogram Options



Red
Channel



Green
Channel




Blue
Channel

3D Histogram
RGB Space

3D array
 $256^3 \sim 16.8M$


22K times more than
3 histograms



Histogram: 1D array of 256 integers
3 colors (RGB): 768 integers


Cumulative Histogram

- Let $h(i)$ be a histogram. $S(i) = \sum_{j=0}^i h(j)$
- $S(0) = h(0)$; $S(255) = \# \text{ of pixels in image}$
- Analogous to CDF – Cumulative Density Function





- What is the area under the curve of h ?
- What is the area under the curve of S ?

Histogram Application: Video Scene Cut Detection






- Similar images inside shots having similar histograms
 - Video Shot Cut Detection: same shot \rightarrow similar histograms
- Compute distance between color histogram of successive frames

Distance Between Histograms (Distributions)






Vector Distance: $|h1-h2| =$
 $= |2-1| + |5-6| + |1-2| + |1-0|$
 $= 4$

But in some cases, this simple vector distance does not work:






- All three have equal vector distance $|h1-h2| = |h1-h3| = 9+9 = 18$
- Seems wrong, since "2" is closer to "1" than "5"

Distance Between Histograms (Distributions)



- A solution: distance between cumulative histograms $S(i) = \sum_{j=0}^i h(j)$




- Vector distance is now OK: $|S1-S2| = 9$


$$|S1-S3| = 36$$

Histogram Equalization Example (Wikipedia)

Original Image
& Histogram





After Histogram
Equalization





Histogram Equalization



- Equal usage of the gray level range. Integration used.



$$P(g) = \int_0^g p(x) dx$$



$$Q(g) = \int_0^g q(x) dx$$



$$g \Rightarrow P^{-1}(Q(g))$$


Histogram Equalization

- Compute Cumulative Histogram $S(i)$ from Histogram $h(i)$

N Pixels, grey levels $0 \dots K$

$h(i) = \#$ pixels at grey level i

$$S(i) = \sum_{j=0}^i h(j)$$



1. Change every original grey level i to $S(i)$
2. Stretch (linear) new gray levels back to $[0 \dots K]$

$$S(m) \rightarrow 0; \quad S(q) \rightarrow K$$


m is the lowest grey level in input image
 q is the highest grey level in input image

$$i \Rightarrow K \frac{S(i) - S(m)}{S(q) - S(m)}$$


Equalization Example

Grey Level (k)	# Pixels (n)	Cumulative Histogram	Scaled 0-7	Round	New Histogram
0	790	790	0.00	0	790
1	1023	1813	2.17	2	0
2	850	2663	3.97	4	1023
3	656	3319	5.35	5	0
4	329	3648	6.05	6	850
5	245	3893	6.57	7	656
6	122	4015	6.83	7	329
7	81	4096	7.00	7	448
Total:	4096				4096

Original Histogram



Equalized Histogram



Histogram Equalization Steps

Target Range is $[0..K]$, can be different from input range

1. Given b/w image $I(x,y)$, create a histogram h :

- For all pixels x,y : $h(I(x,y)) = h(I(x,y)) + 1$

2. [Hist] Create cumulative histogram $S(i)$:

- $S(0) = h(0)$; $S(i+1) = S(i) + h(i+1)$;

Let \underline{m} and \underline{q} be the smallest & highest input grey levels

1. [Hist] Create Look Up Table (LUT) $T(i)$:


- $T(i) = \text{round} \{ K \times [S(i) - S(\underline{m})] / [S(\underline{q}) - S(\underline{m})] \}$

2. Apply LUT T to image I , get equalized image J :


- $J(x,y) = T(I(x,y))$

Example of Equalization (Wikipedia)

Original Image
& Histogram

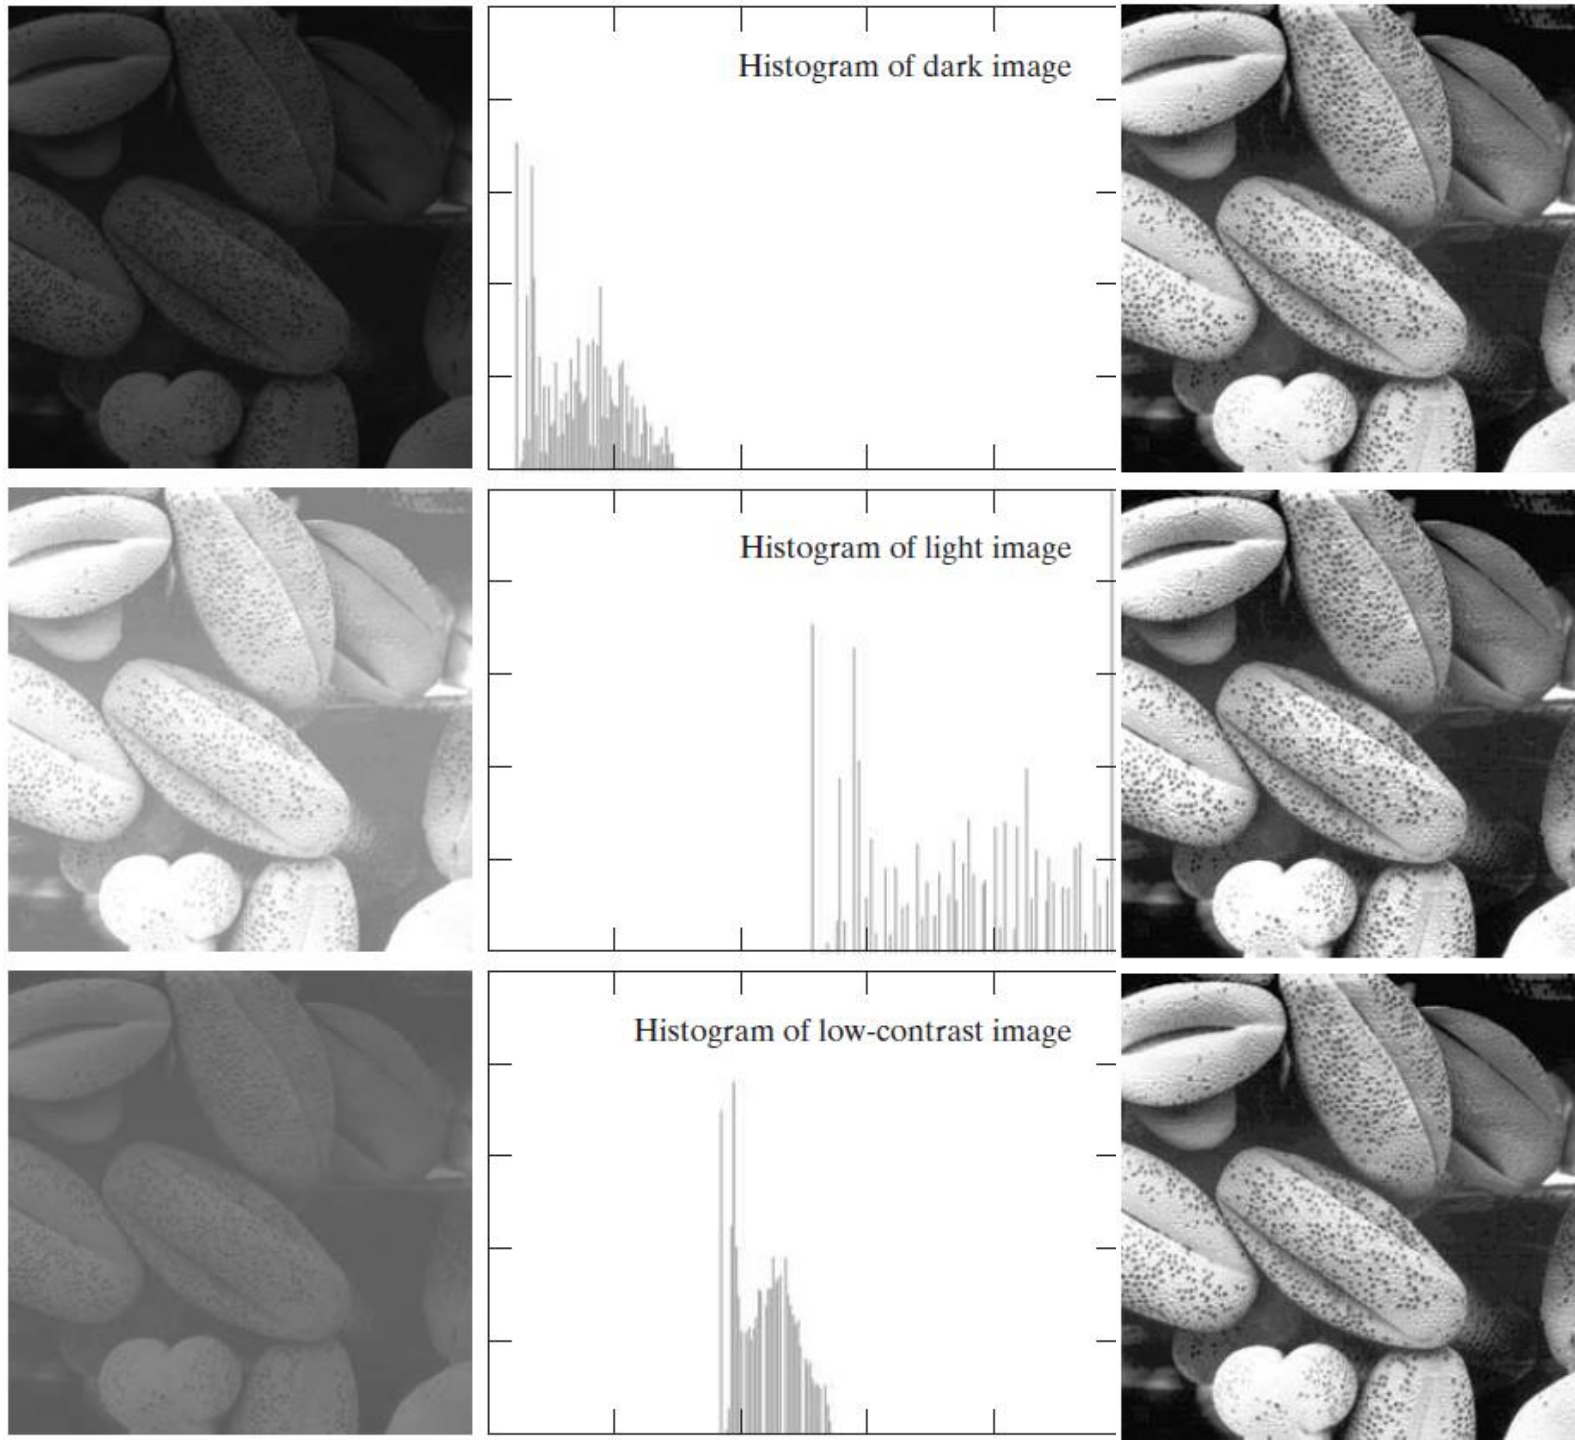


After Histogram
Equalization



Properties of Histogram Equalization


- Monotonic Transformation: Does not reverse intensity order
- What happens if we apply Equalization twice?
- New intensity \approx Cumulative Probability
 - What is the meaning of grey level 127 in an image after equalization to [0..255]?
- Will fail(?) if assumptions are not true.
 - When?




Histogram
Equalization is
invariant to
monotonic point
operations

Histogram Specification


Histogram of Input Image




Specified Histogram



Cumulative Histogram
of Input Image $f_1[n, m]$



Cumulative Histogram
of Specified Histogram




- Change the histogram of input image to any specified histogram
- Input histogram H_1 and cumulative C_1
- Target histogram H_0 and cumulative C_0
- For each input grey level u find the target grey level v such that $C_1(u) = C_0(v)$

Adaptive Histogram Equalization

- Different intensity distributions in an image
 - Example: sunny areas and shadowed areas
- Poor result for Histogram Equalization
 - Work on sunny and shadowed areas separately
 - How to segment?
- Workaround: Compute histogram in local regions around each pixel

Adaptive Histogram Equalization



- For each pixel
 - Compute Equalization LUT in local region
 - Transform by LUT only the center pixel
- Go to next pixel
- How to optimize?

Adaptive Histogram Equalization

Original



Global Equalization



Window = 100x100



Window = 50x50



Quantization: Reduce # of Colors

- Example: To reduce $[0..255]$ to 32 grey levels $[0..31]$, we use a Quantization LUT

In	0	1	...	k	...	255
Out	q_0	q_0	...	q_n	...	q_{31}


- Uniform Quantization: Every grey level k is mapped to $k/8$
- To restore the original image, we use a Restoration LUT

In	0	1	...	k	...	31
Out	...	G_1	...	G_k	...	G_{31}

- Example: Every q is mapped to $q \times 8$ (?)
- When is uniform Quantization not optimal?


Uniform Quantization

- E.g. - Every grey level k is mapped to $k/8$
- To restore the original image every q is mapped to $q \times 8$ (?)



0 .. 31	→ 0	→ 0
32 .. 63	→ 32	→ 36
64 .. 95	→ 64	→ 72
96..127	→ 96	→ 108
128..159	→ 128	→ 142
...		...
224..255	→ 224	→ 255

Quantization



- Divide grey level range into fewer segments
- Quantization: A grey level segment is mapped to one index
 - Borders of segments: $z_0=0, z_1, z_2, \dots, z_k=255$
 - All grey levels in segment $[z_{i-1}, z_i]$ are mapped to $i-1$
- Restoration: Each index i will be restored to intensity q_i
- Uniform Quantization:

$$z_{i+1} - z_i = (z_k - z_0) / k$$

$$q_i = (z_i + z_{i+1}) / 2$$

Quantization Error

- Assume that during quantization grey level g is coded by i , and code i is restored as q_i
- A Possible quantization error **for one pixel p** is:

$$E_p^2 = (g - q_i)^2$$

- The total error introduced by quantization of all pixels is:

$$E^2 = \sum_{pixels\ p} E_p^2 = \sum_{i=0}^{255} hist(i) E_i^2$$


All pixels with same grey level will have same quantization error

- Unknown optimal transformations: $g \rightarrow i, i \rightarrow q_i$

SSD As a Quantization Error

- SSD (L_2 , Sum of Squared Differences) is a poor error measure compared to human perception
- There are suggestions for other measures, but they are harder to compute and to analyze (E.g. L_1 , Sum of Absolute Values)
- $SSD = 0$ implies the same image...
- (Best way to compute image similarity is with Neural networks...)

Optimal Quantization - Condition



- Minimize the error:
- Solution (**Prove!**):

$$\sum_{i=0}^{k-1} \left(\sum_{g=\lfloor z_i \rfloor + 1}^{\lfloor z_{i+1} \rfloor} (q_i - g)^2 h(g) \right)$$

$$q_i = \frac{\sum_{g=\lfloor z_i \rfloor + 1}^{\lfloor z_{i+1} \rfloor} g \cdot h(g)}{\sum_{g=\lfloor z_i \rfloor + 1}^{\lfloor z_{i+1} \rfloor} h(g)}$$

$$z_i = \frac{q_{i-1} + q_i}{2}$$

Optimal Quantization - Process

- Find z_i and q_i such that

$$q_i = \frac{\sum_{g=\lfloor z_i \rfloor + 1}^{\lfloor z_{i+1} \rfloor} g \cdot h(g)}{\sum_{g=\lfloor z_i \rfloor + 1}^{\lfloor z_{i+1} \rfloor} h(g)} \quad z_i = \frac{q_{i-1} + q_i}{2}$$

- This is done iteratively
- Initial guess: find z_i such that for all i there are same number of pixels whose grey level is between z_i and z_{i+1} .
- From initial guess of z_i compute q_i .
- Iterate: computing z_i and then q_i until convergence.

Next: Fourier

- Discrete Fourier Transform in Wikipedia
- <http://homepages.inf.ed.ac.uk/rbf/HIPR2/fourier.htm>
- <https://betterexplained.com/articles/an-interactive-guide-to-the-fourier-transform/>