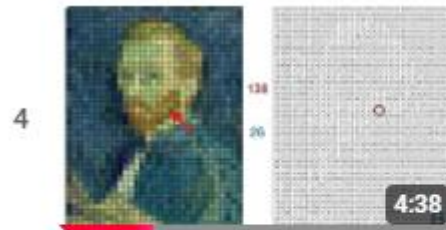


Videos to Watch (by Steve Seitz)

<https://www.youtube.com/playlist?list=PLWfDJ5nla8UpwShx-lzLJqcp575fKpsSO>



Images in 5 minutes: The Case of the Splotched Van Gogh, Part 1

Graphics in 5 Minutes • 4.7K views • 2 years ago

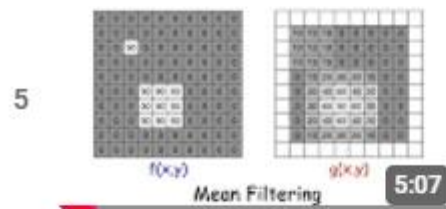


Image filtering in 5 minutes: The Case of the Splotched Van Gogh, Part 2

Graphics in 5 Minutes • 5.3K views • 2 years ago



Fourier Transform in 5 minutes: The Case of the Splotched Van Gogh, Part 3

Graphics in 5 Minutes • 6.8K views • 2 years ago

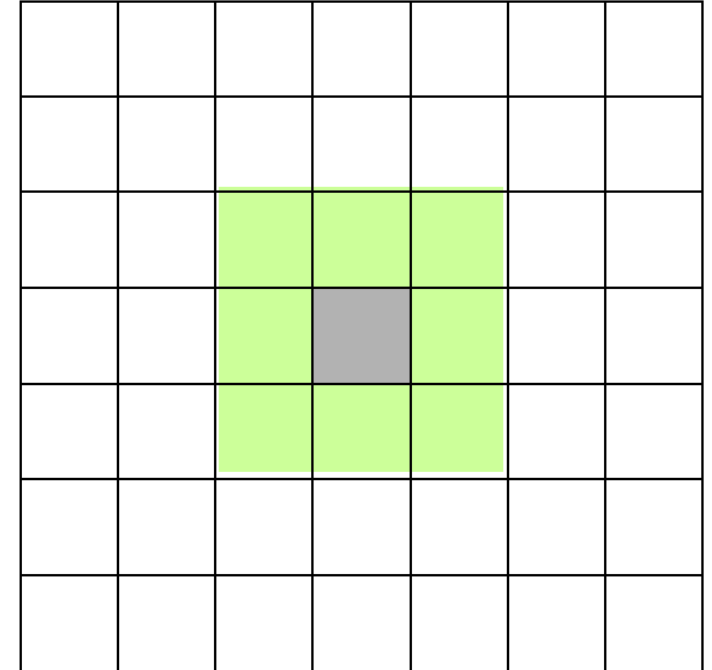
- 4) Sampling, Moire
- 5) Convolutions, Blur before Sample
- 6) Fourier – Explaining the sampling effect

Convolution

- A linear operator (weighted sum of neighbors) applied identically on all pixels.
- Example - **Blur**: Average a pixel with its 3×3 neighborhood

$$h(x, y) = \sum_{i=-1}^1 \sum_{j=-1}^1 f(i, j) g(x - i, y - j)$$

Blurred Pixels Kernel Original Pixels



- Image g is blurred by kernel f giving image h (uniform average)

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

(x,y) of center = (0,0)

Convolution

$$h(x, y) = \sum_{i=-1}^1 \sum_{j=-1}^1 f(i, j) g(x - i, y - j)$$

- A linear operator (weighted sum of neighbors) applied identically on all pixels.
 - Blur: Average a pixel with its neighbors (weighted average)
 - Why is blur important?

$\frac{1}{16}$	1	2	1
	2	4	2
	1	2	1

(x,y) of center = (0,0)

0	0	0
0	1	0
0	0	0

- Do Nothing kernel

Convolution

$$h(x, y) = \sum_{i=-1}^1 \sum_{j=-1}^1 f(i, j) g(x - i, y - j)$$

- Shift the image to the left (reflection of kernel)

Only a single non-zero

$$h(x, y) = f(-1, 0) g(x+1, y)$$

1		

(x,y) of center = (0,0)

- Edge: Compute a difference between neighbors of a pixel (~derivative)

- Vertical edges

1	-1
---	----

1/2		-1/2

- Horizontal edges

1
-1

	1/2	
	-1/2	

- 1D convolution, 2D convolution

Convolution - Blur

- Blur: Replace a pixel with an average of its neighbors

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

- Horizontal blur

1/3	1/3	1/3

- Vertical Blur

	1/3	
	1/3	
	1/3	

- Diagonal Blur

		1/3
	1/3	
1/3		

Convolutions: Smoothing, Noise Cleaning

Smoothing
($\sum \text{weights} = 1$)

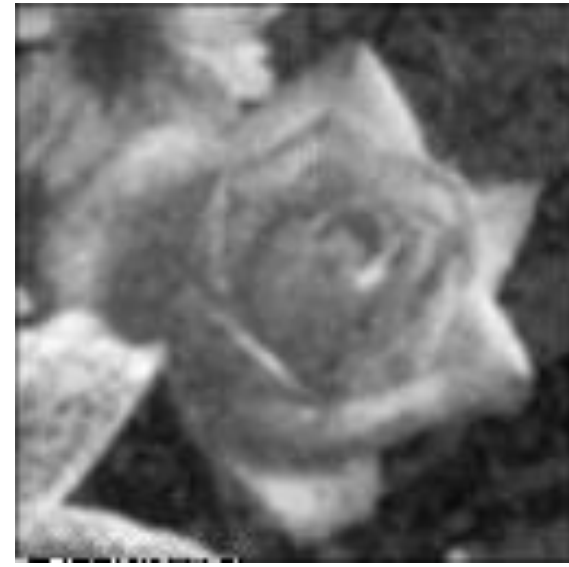
$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$



Original Image



Noisy Image



Filtered Image

1-D Discrete Convolution (*)

$f = (f_0, f_1, f_2, f_3, \dots)$, $g = (g_0, g_1, g_2, g_3, \dots)$, are 1-D arrays

$$h = f * g; \quad h(x) = (f * g)(x) = \sum_{a=1}^n f(a)g(x-a)$$

Let f, g be defined in the coordinate range $[0..5]$. Assume cyclic boundaries.

$$f = (0,0,0,1,0,0)$$

$$g = (0,0,0,1,-1,0)$$

$$h(6) = h(0) = \sum_{a=0}^5 f(a) \cdot g(6-a) = f(3) \cdot g(3) = 1$$

$$h(7) = h(1) = \sum_{a=0}^5 f(a) \cdot g(7-a) = f(3) \cdot g(4) = -1$$

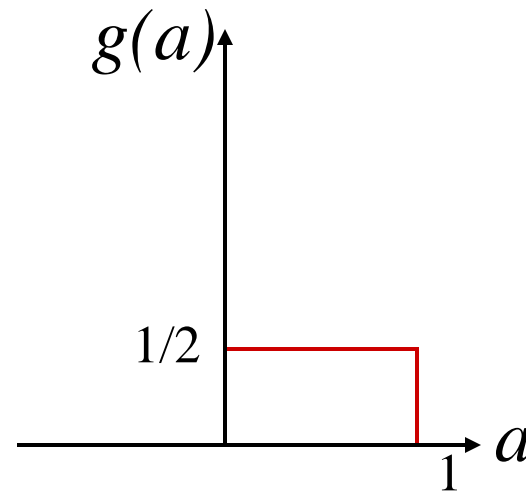
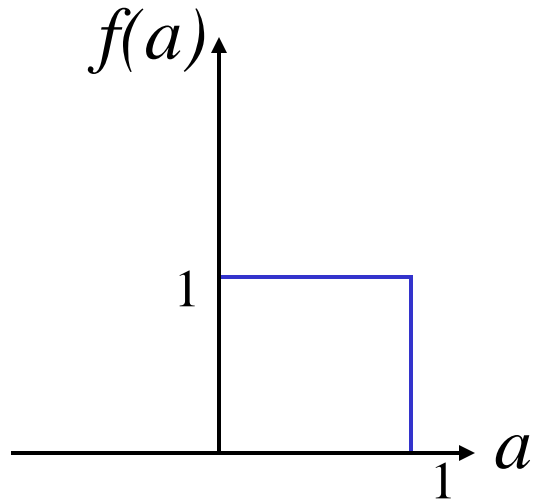
Boundary Handling

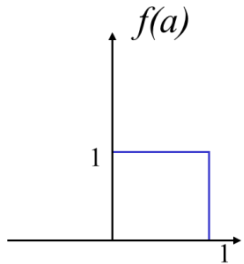
$$f = (0, 0, 0, 1, 0, 0)$$

- **Q:** f is defined in the range of $[0..5]$. What is $f(-3)$; $f(9)$?
- **A1:** $f(-3) = f(9) = 1$; Cyclic approach (Fourier). $-3 = 9 = 3 \text{ (Mod. 6)}$
[Rarely Used in practice]
- **A2:** $f(-3) = f(9) = 0$; Every index out of range is zero.
[Zero Padding]
- **A3:** Reflection. $F(-a) = F(a)$. $F(N+a) = F(N-1-a)$;

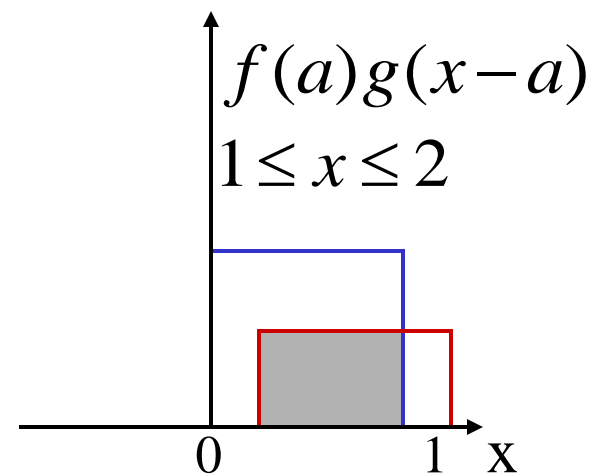
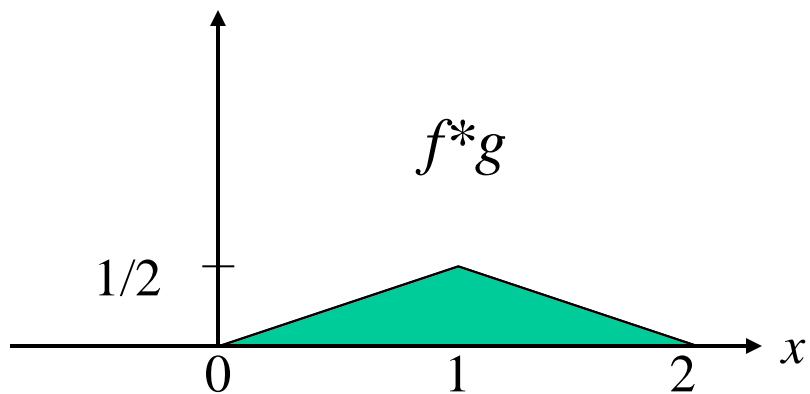
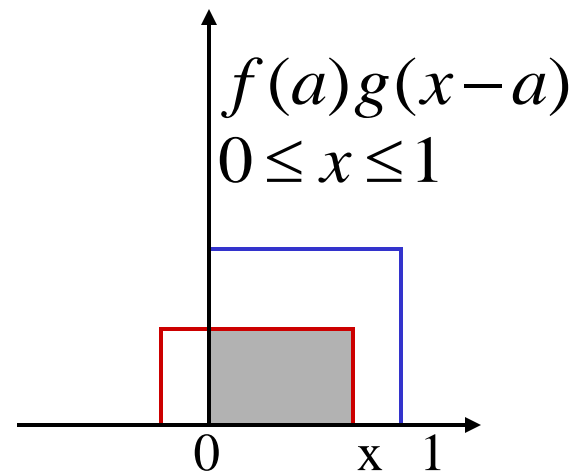
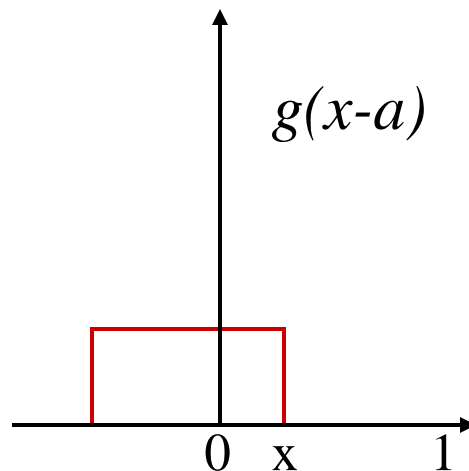
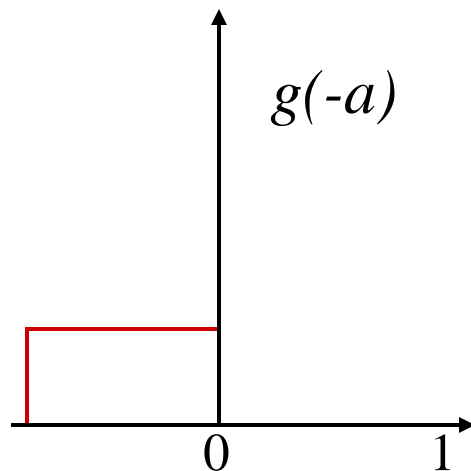
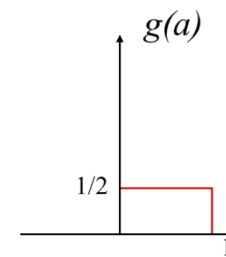
1-D Continuous Convolution

$$(f * g)(x) = \int_{-\infty}^{\infty} f(a)g(x-a)da$$



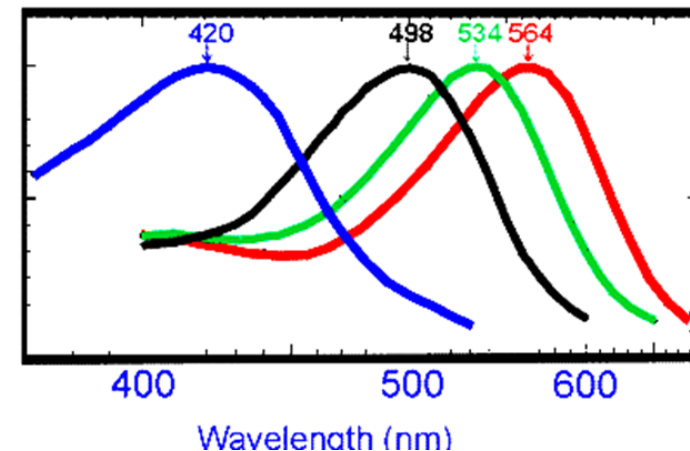
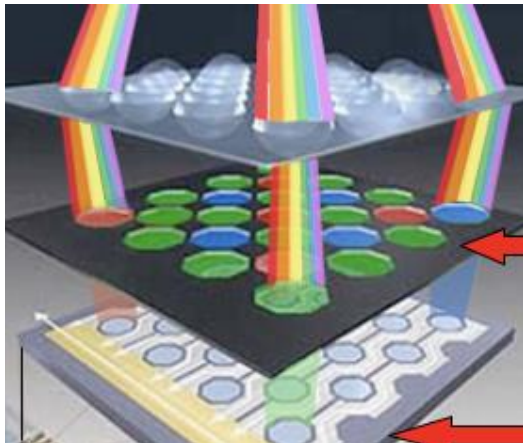


$$(f * g)(x) = \int_{-\infty}^{\infty} f(a)g(x-a)da$$



Why Continuous Convolution?

- The physical world is continuous (sensors)
 - Response of sensors to color (color space – electromagnetic wavelength)
 - Area response of sensors (image space)



Convolution Questions

(i) What other courses use convolutions?

Algorithms (polynom, FFT), Deep Learning (CNN)

(i) What is the complexity of discrete convolution?

$O(N^2)$; With FFT it is $O(N \log N)$

(i) Why “reflect”?

Many good properties, e.g. Commutative

Convolution Theorem (Φ is Transform Fourier)

$$\Phi(f * g) = F \cdot G \quad \text{Pointwise Multiplication}$$

$$\Phi(f \cdot g) = F * G$$

Convolution in spatial domain $(f(x,y), g(x,y))$ is equivalent to **pointwise multiplication** in frequency domain $(F(u,v), G(u,v))$

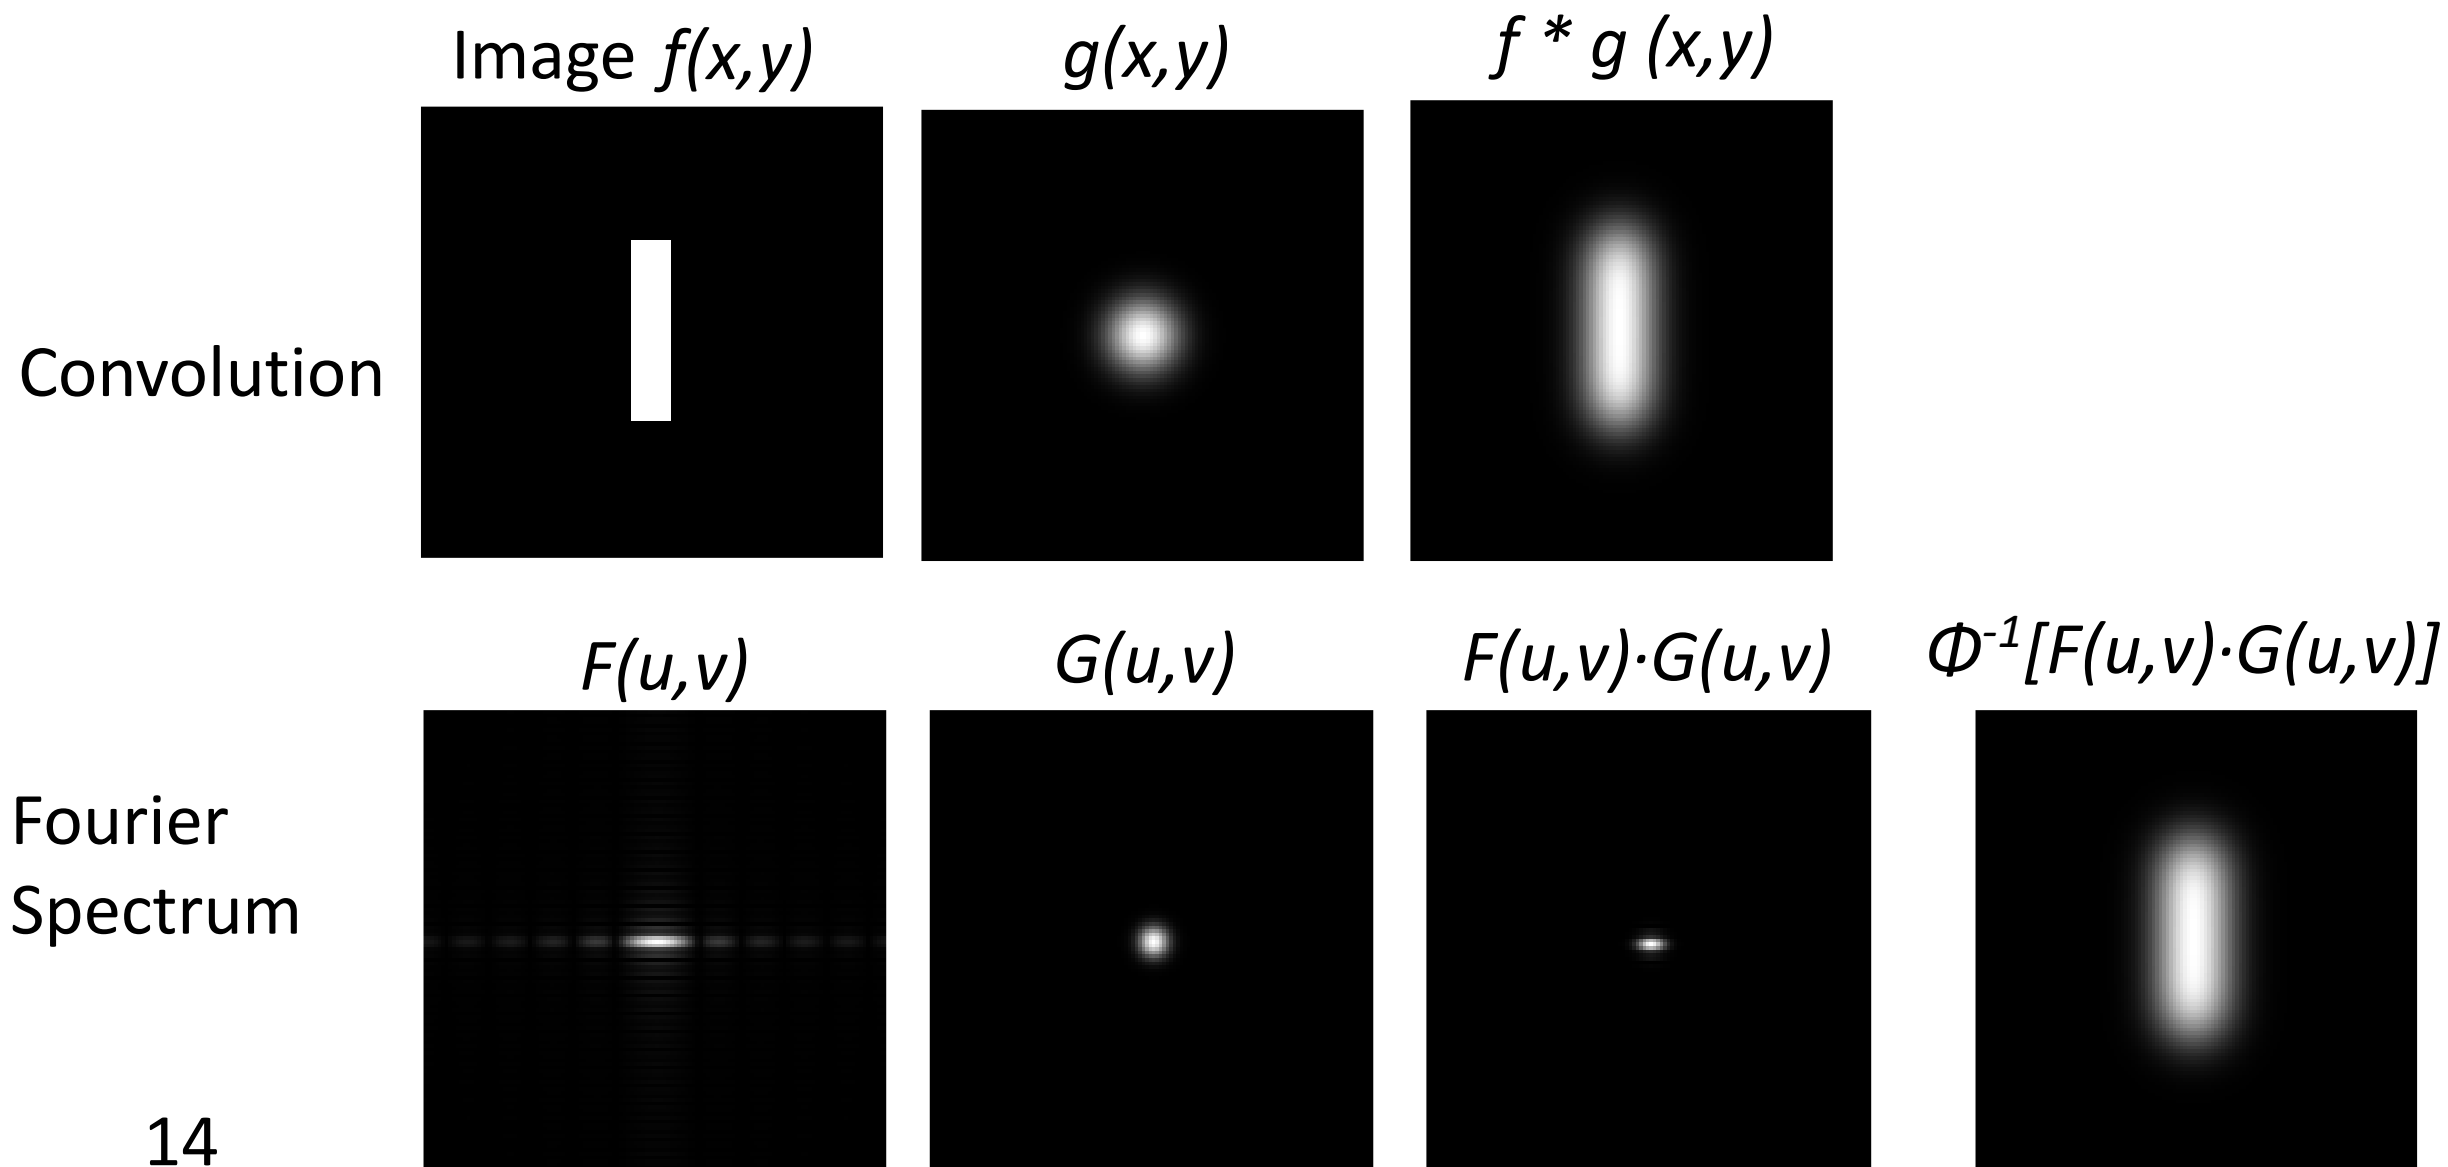
Performing Convolution by using Fourier:

$$f * g = \Phi^{-1}(F \cdot G) = \Phi^{-1}(\Phi(f) \cdot \Phi(g))$$

FFT (Fast Fourier Transform) reduces complexity of convolution:

$$O(N^2) \rightarrow O(N \log N)$$

Convolutions in the Frequency Domain



Properties of Convolution

Commutative: $f * g = g * f$

Associative : $f * (g * h) = (f * g) * h$

Distributive: $f * (g + h) = f * g + f * h$

Convolution is a Linear Operation over 1D vectors and 2D images

Convolution as Matrix Multiplication (Cyclic)

A Linear Operator can be expressed as matrix multiplication

$$(x_1, x_2, x_3, x_4, x_5, x_6) * \frac{1}{4} \begin{pmatrix} 1 & 2 & 1 \end{pmatrix}$$

$$(x_1, x_2, x_3, x_4, x_5, x_6) \cdot \frac{1}{4} \begin{pmatrix} 2 & 1 & 0 & 0 & 0 & 1 \\ 1 & 2 & 1 & 0 & 0 & 0 \\ 0 & 1 & 2 & 1 & 0 & 0 \\ 0 & 0 & 1 & 2 & 1 & 0 \\ 0 & 0 & 0 & 1 & 2 & 1 \\ 1 & 0 & 0 & 0 & 1 & 2 \end{pmatrix}$$

Circulant Matrix

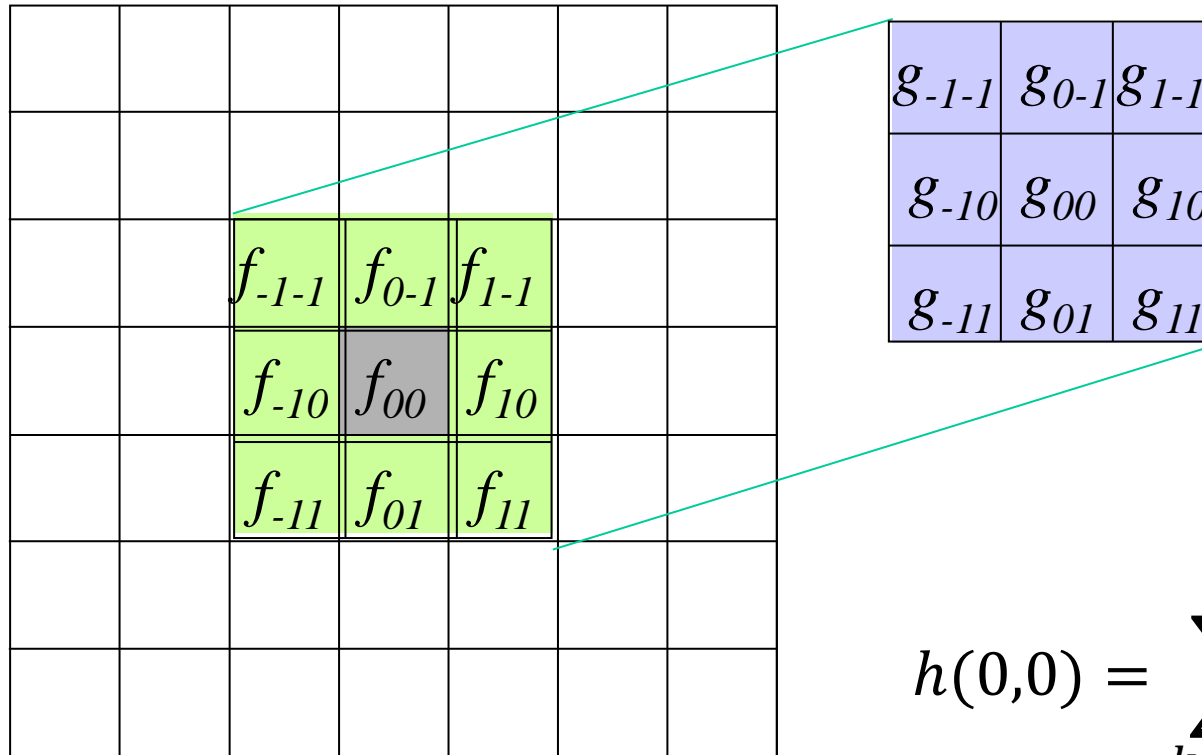
2-D Discrete Convolution

$$h = f * g$$

$$h(x, y) = \sum_{k=0}^n \sum_{l=0}^m f(k, l) g(x - k, y - l)$$

Blurred Pixel Kernel Original Pixels

2D Discrete Convolution



$$h(0,0) = \sum_{k=-1}^1 \sum_{l=-1}^1 f(k,l) \cdot g(0-k, 0-l)$$

$$f(-1,-1)g(1,1) + f(-1,0)g(1,0) + f(-1,1)g(-1,1) + f(0,-1)g(0,1) + \dots$$

Simple Convolutions (Empty is zero)

$$\begin{bmatrix} 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

Simple Convolutions (Empty is zero)

$$[1 \quad 1] * [1 \quad 1] = [1 \quad 2 \quad 1]$$

$$\begin{bmatrix} 1 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix}$$

$$[1 \quad 1] * \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$$

$$[1 \quad 2 \quad 1] * \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

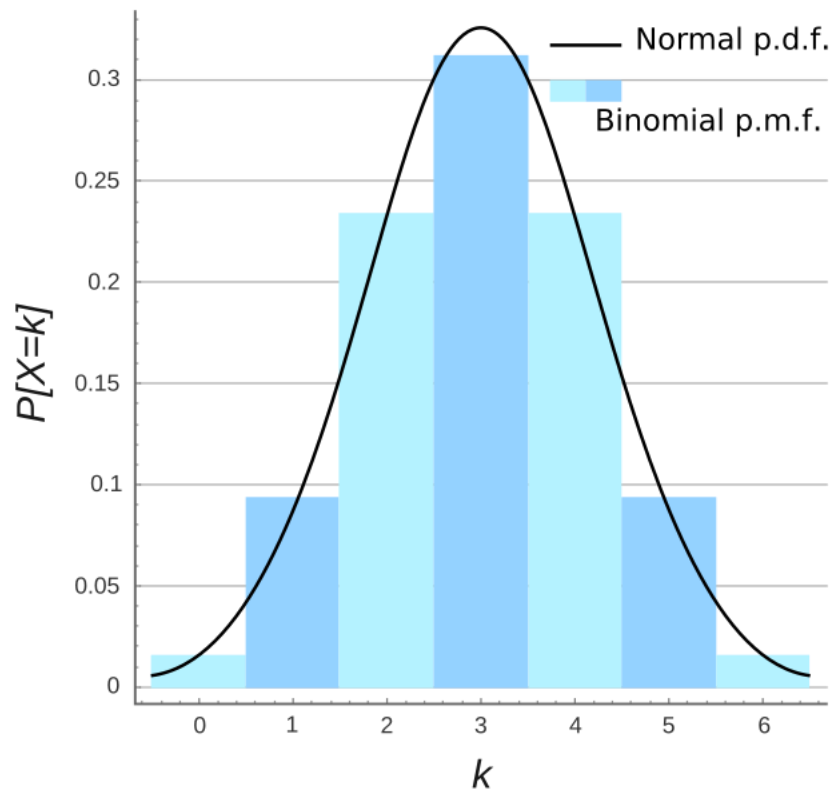
$$\begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

$$[1 \quad 1] * [1 \quad 1] * \begin{bmatrix} 1 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 \\ 1 \end{bmatrix} = [1 \quad 1] * \begin{bmatrix} 1 \\ 1 \end{bmatrix} * [1 \quad 1] * \begin{bmatrix} 1 \\ 1 \end{bmatrix} = [1 \quad 1] * \begin{bmatrix} 1 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 \\ 1 \end{bmatrix} * [1 \quad 1]$$

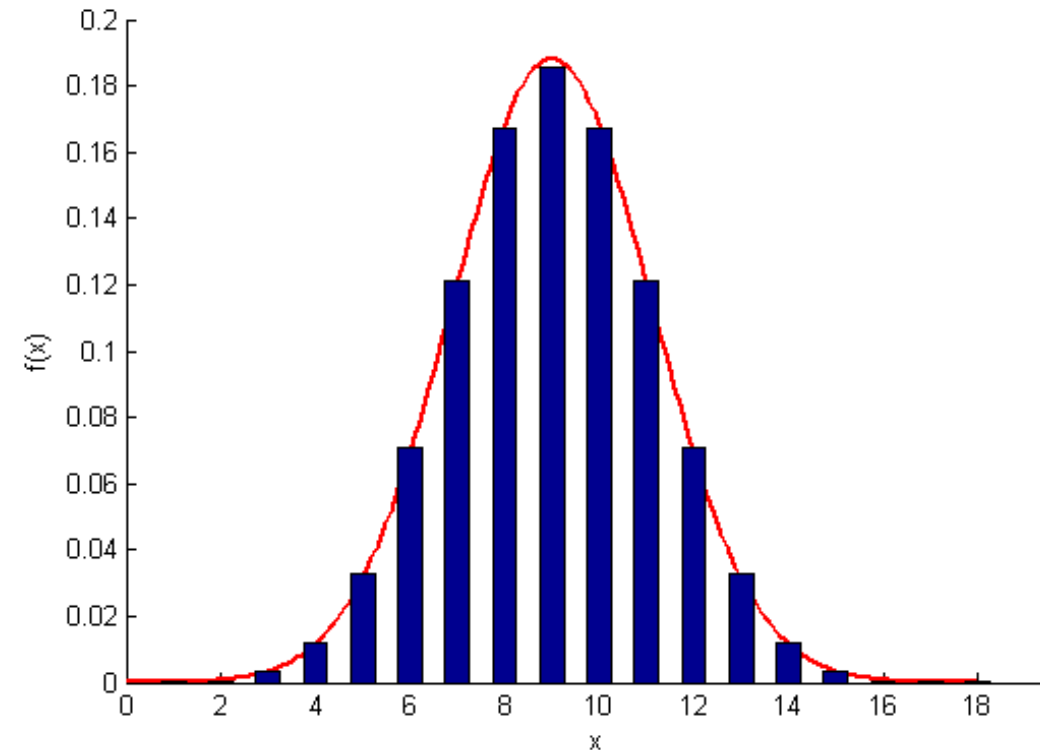
Simple Convolutions (Empty is zero)

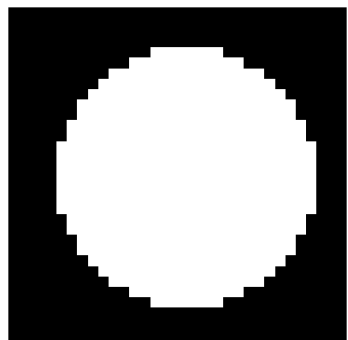
- Repetitive convolutions of [1 1] giving binomial coefficients (Pascal triangle)

- For large n , binomial coefficients approximate a Gaussian $G(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$



$$\begin{array}{c} [1 \ 1]^n \\ [1 \ 1] \\ [1 \ 2 \ 1] \\ [1 \ 3 \ 3 \ 1] \\ [1 \ 4 \ 6 \ 4 \ 1] \end{array}$$

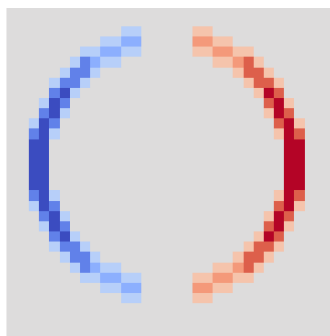




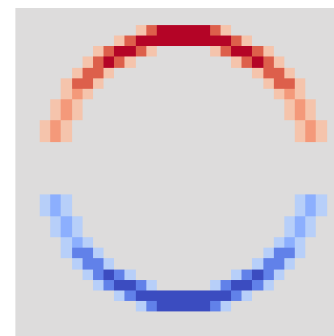
Edge Detection

Edge: Large difference between neighboring pixels

- Vertical edges
- right-left diff



- Horizontal edges
- below-above diff



$$\begin{bmatrix} 1/2 & 0 & -1/2 \end{bmatrix} * \begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix} = \begin{bmatrix} 1/6 & 0 & -1/6 \\ 1/6 & 0 & -1/6 \\ 1/6 & 0 & -1/6 \end{bmatrix}$$

$$\begin{bmatrix} 1/2 \\ 0 \\ -1/2 \end{bmatrix} * \begin{bmatrix} 1/3 & 1/3 & 1/3 \end{bmatrix} = \begin{bmatrix} 1/6 & 1/6 & 1/6 \\ 0 & 0 & 0 \\ -1/6 & -1/6 & -1/6 \end{bmatrix}$$

- Difference, a derivative, high pass - across the desired edge
- Blur, low pass, noise cleaning - along the desired edge

Convolutions: Smoothing, Noise Cleaning

Q: What is the average gray level after smoothing?

Smoothing
($\sum \text{weights} = 1$)

$$\frac{1}{5} \begin{bmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

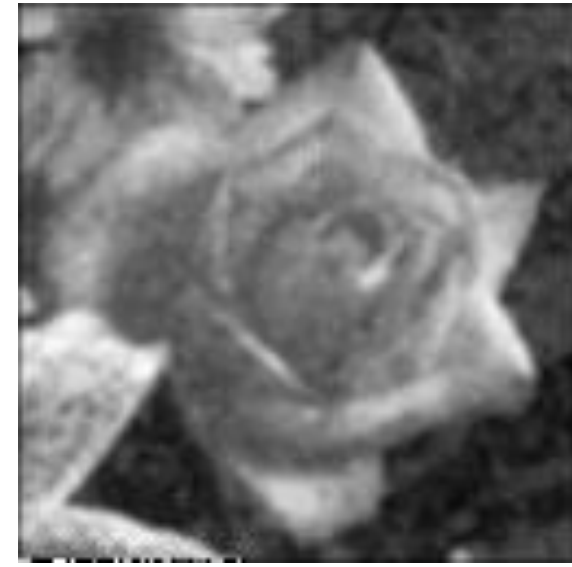
$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$



Original Image



Noisy Image



Filtered Image

Noise Assumption: Additive, Zero-Mean, Independent

Convolutions: Smoothing, Noise Cleaning

Original Picture



Convolutions: Smoothing, Noise Cleaning

Added Zero-Mean Noise

Grey=0;

Black=-255; White=255

- Noise is random for every pixel, and does not depend on the noise at any other pixel.
- White Noise: The Fourier spectrum of White Noise is flat – All frequencies have same amplitude.

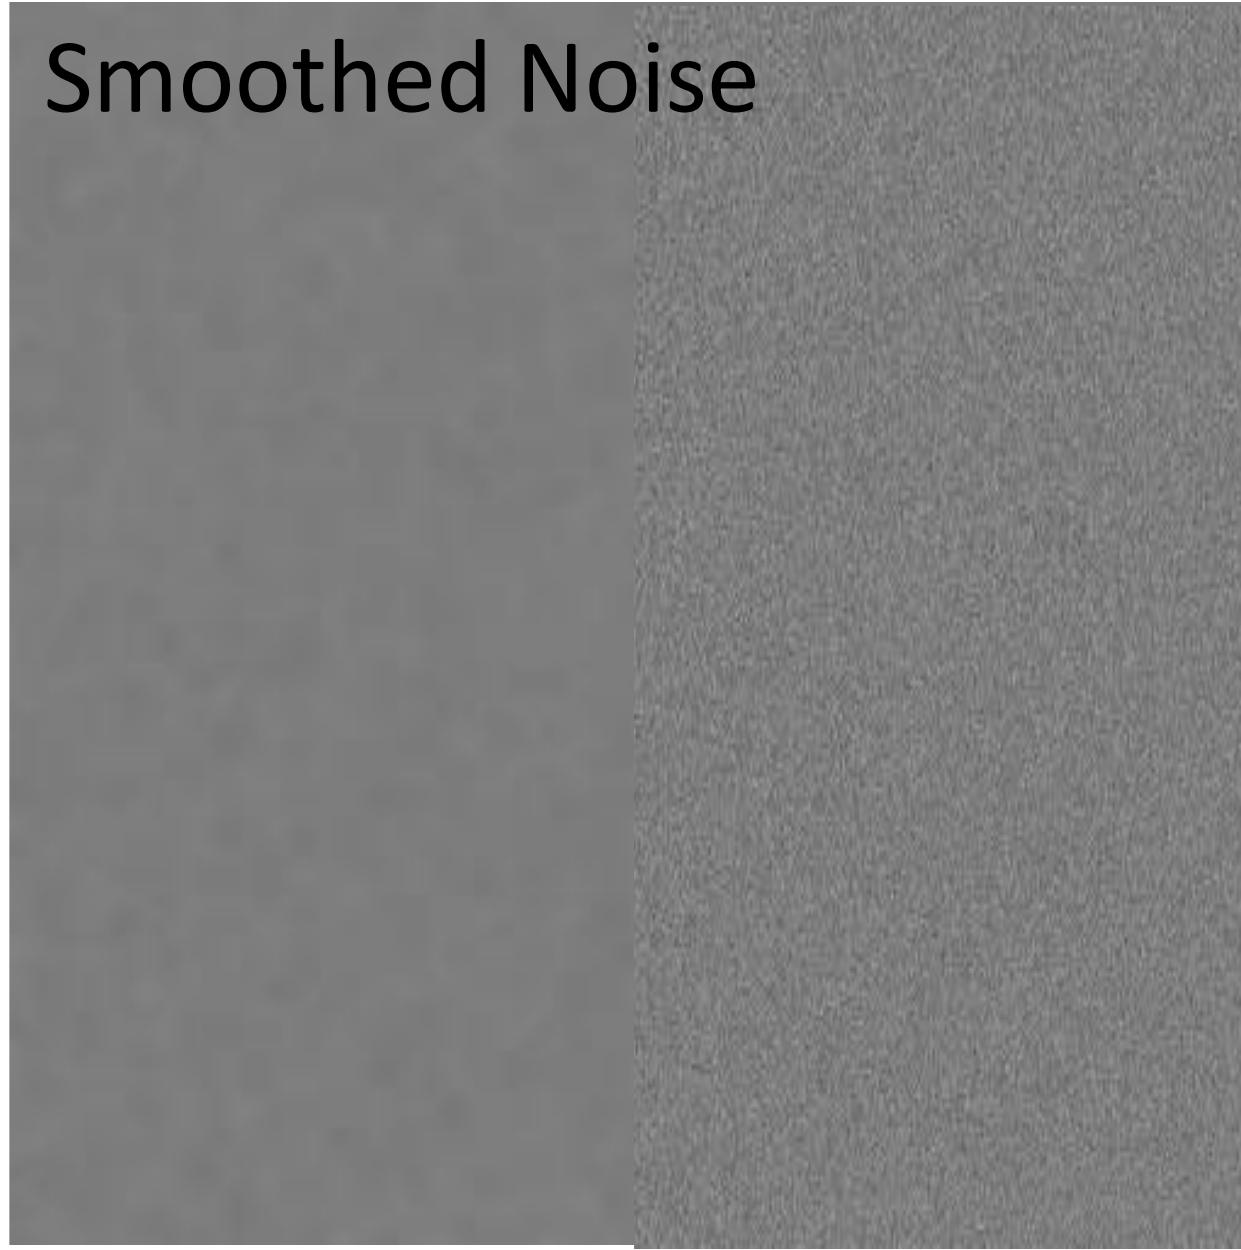
Convolutions: Smoothing, Noise Cleaning

Noisy Picture



Convolutions: Smoothing, Noise Cleaning

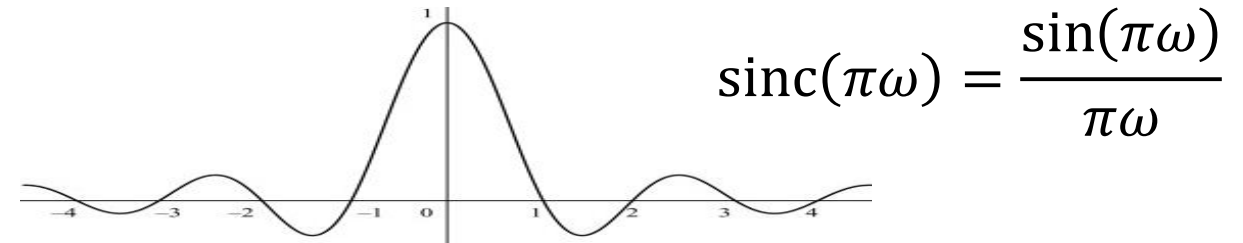
Smoothed Noise



Low Pass: Fourier Domain & Image Domain

- Image: convolve f by box Fourier: Multiply F by Sinc

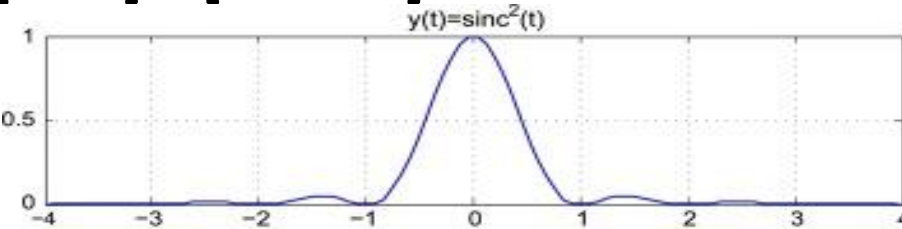
- $[1 \ 1]$



- Multiply Fourier F by box \Leftrightarrow Convolve image f with Sinc

- Image: convolve by $[1 \ 1] * [1 \ 1] = [1 \ 2 \ 1]$

- Fourier: Multiply F by Sinc^2



- Image: Gaussian \Leftrightarrow Fourier: Gaussian

- Blur image f w. Gaussian \Leftrightarrow Multiply Fourier F w. Gaussian

Convolution: Spatial filtering



Smoothing by Convolution


$$\frac{1}{9} \times$$

1	1	1
1	1	1
1	1	1

$$\frac{1}{16} \times$$

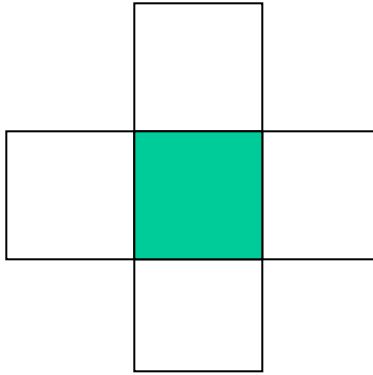
1	2	1
2	4	2
1	2	1

Smoothing by Spatial Filtering

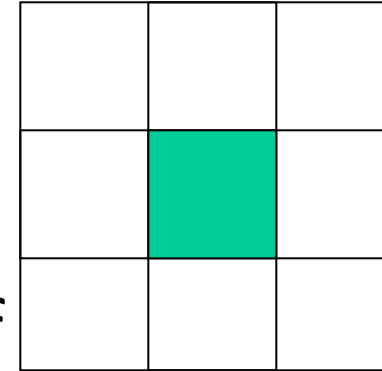


Using larger filter

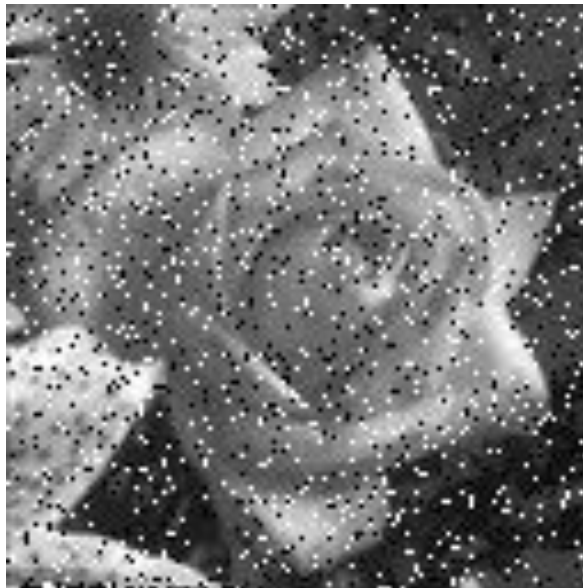
Noise Cleaning by Median Filtering



- Replace the value of a pixel with the MEDIAN of its neighborhood
- Depends on the definition of “neighborhood”



Original



Salt & Pepper Noise



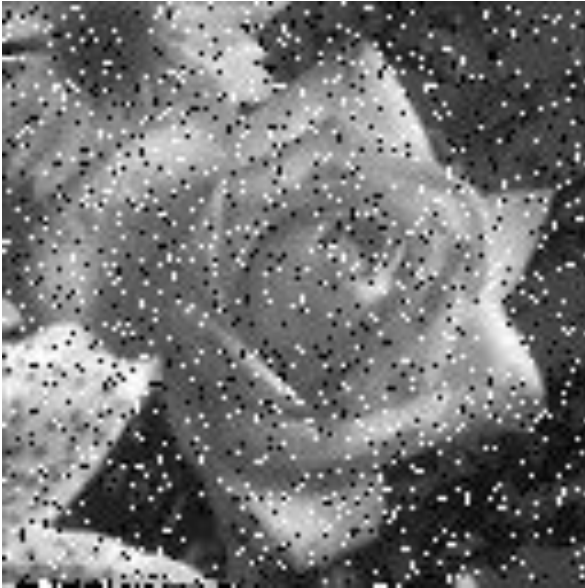
Median

Median

- Given a set of numbers, **half** are above the median and **half** are below the median.
- 1, 3, 3, **6**, 7, 8, 9 Median = 6 [Mean: 5.2]
- 1, 2, 3, **4, 5**, 6, 8, 9 Median = 4.5 [Mean: 4.75]
- 1, 2, 1, 2, 10^6 Mean = $(1,000,006)/7 = 200,001$
- 1, 1, **2**, 2, 10^6 Median = 2
- Median is **robust** to outliers

Noise Cleaning

- Averaging / Smoothing - Loss of Detail
- Median - Blockiness

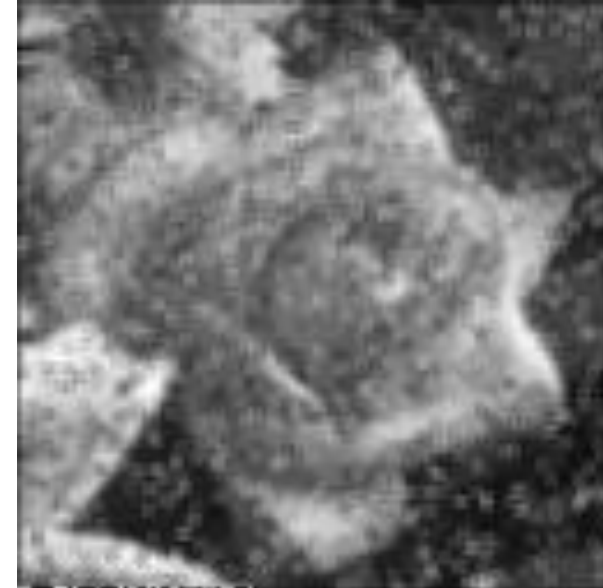


Salt & Pepper Noise

Black or white, ignores original value



Median



Blur (Convolution)

Spatial Approximation to Derivatives

$$\frac{\partial}{\partial x} f(i, j) = \lim_{h \rightarrow 0} \frac{f(i, j) - f(i - h, j)}{h} \cong f(i, j) - f(i - 1, j) \quad \text{convolution with } \begin{pmatrix} 1 & -1 \end{pmatrix}$$

$$\cong \frac{f(i + 1, j) - f(i - 1, j)}{2} \quad \frac{1}{2} \begin{pmatrix} 1 & 0 & -1 \end{pmatrix}$$



$\begin{pmatrix} 1 & -1 \end{pmatrix}$ Derivative between 2 pixels

$\frac{1}{2} \begin{pmatrix} 1 & 0 & -1 \end{pmatrix}$ Derivative at center pixel

Spatial Approximation to Derivatives

$$\frac{\partial}{\partial x} f(i, j) = \lim_{h \rightarrow 0} \frac{f(i, j) - f(i - h, j)}{h} \cong f(i, j) - f(i - 1, j)$$

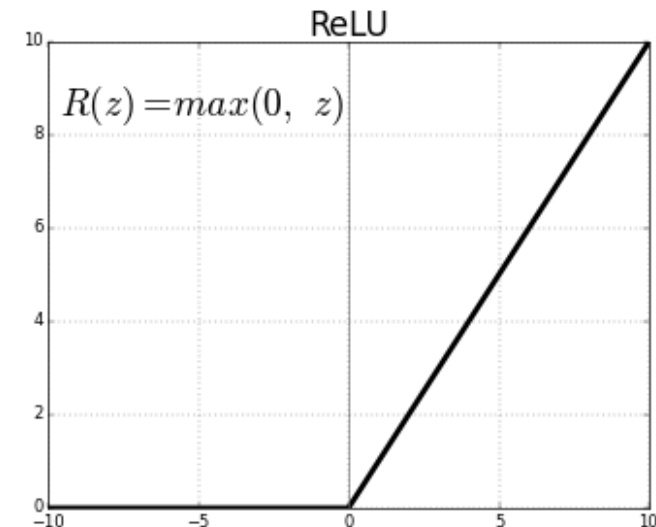
$$\cong \frac{f(i + 1, j) - f(i - 1, j)}{2}$$

convolution with $\begin{pmatrix} 1 & -1 \end{pmatrix}$

Convolution with $\frac{1}{2} \begin{pmatrix} 1 & 0 & -1 \end{pmatrix}$



Detection of Vertical Edges (ReLU: Negative \rightarrow 0)



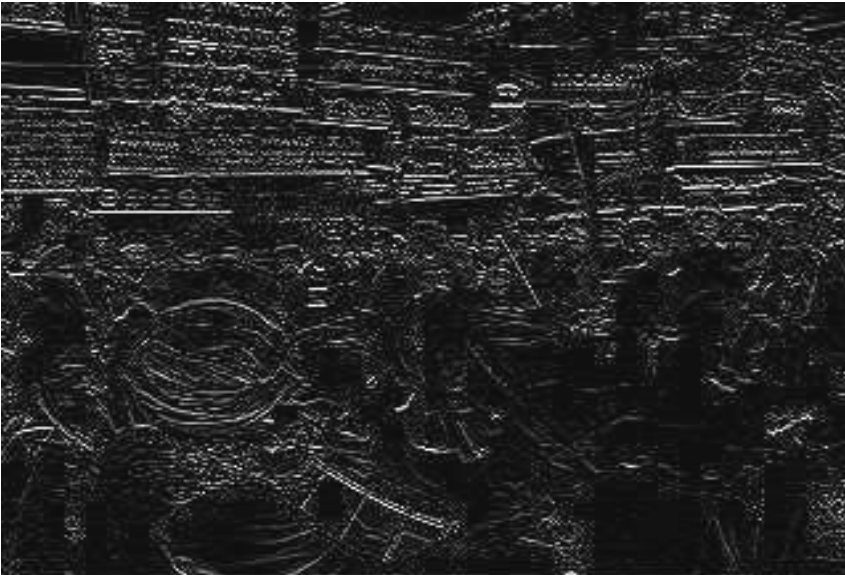
Spatial Approximation to Derivatives

Q: What is the average derivative in all pixels?

$$\frac{\partial}{\partial y} f(i, j) \cong f(i, j) - f(i, j - 1)$$

convolution with $\begin{pmatrix} 1 \\ -1 \end{pmatrix}$

or with $\frac{1}{2} \begin{pmatrix} 1 \\ 0 \\ -1 \end{pmatrix}$

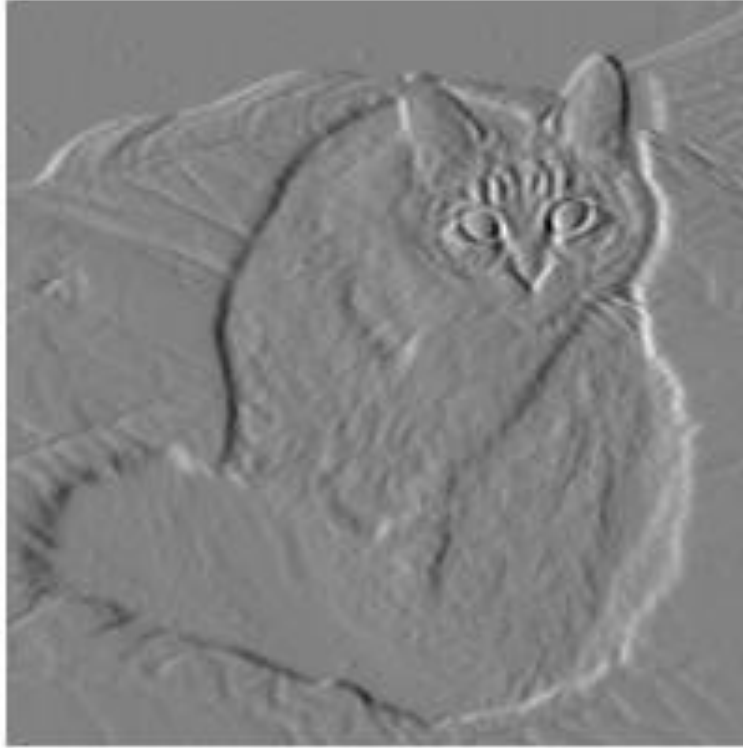


Horizontal Edge Detection (ReLU)

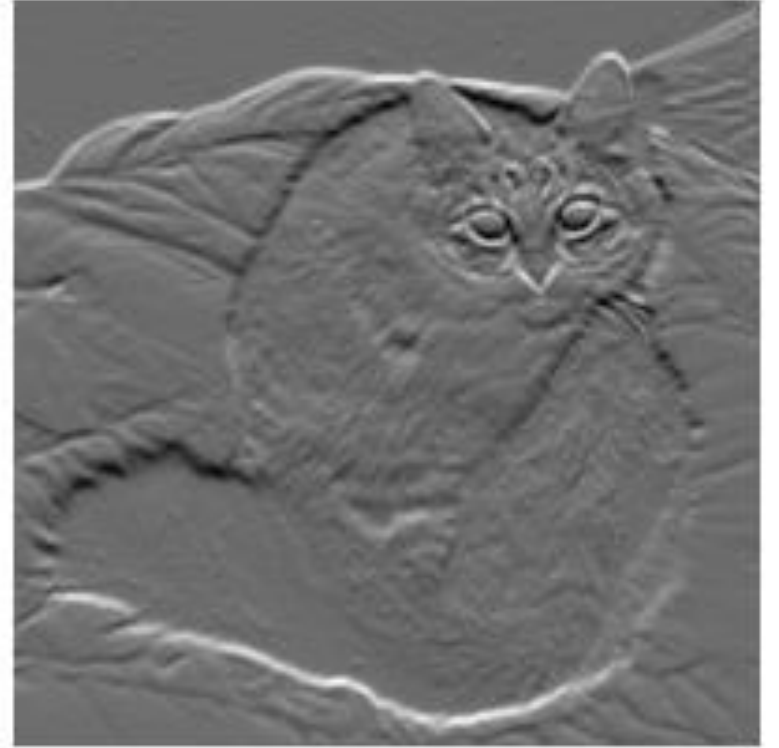
Derivative Example



Image



x derivative



y derivative

Negative: Black. Zero: grey. Positive: White.

No absolute value. No ReLU

Image Derivatives (cont')

Since a picture is not continuous, there are many **approximations** to the derivative

Popular blur kernels (Sobel):

$$\frac{\partial f}{\partial x} = \frac{1}{8} \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$\frac{\partial f}{\partial y} = \frac{1}{8} \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Note: Good derivative filters are edge in one direction and **blur** in the orthogonal direction.

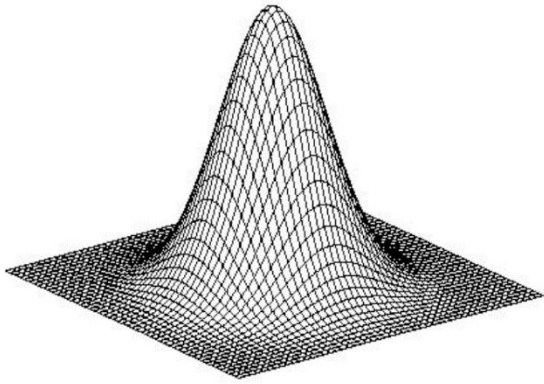
Compare to:

$$\frac{\partial f}{\partial x} = \begin{bmatrix} 1 & -1 \end{bmatrix}$$

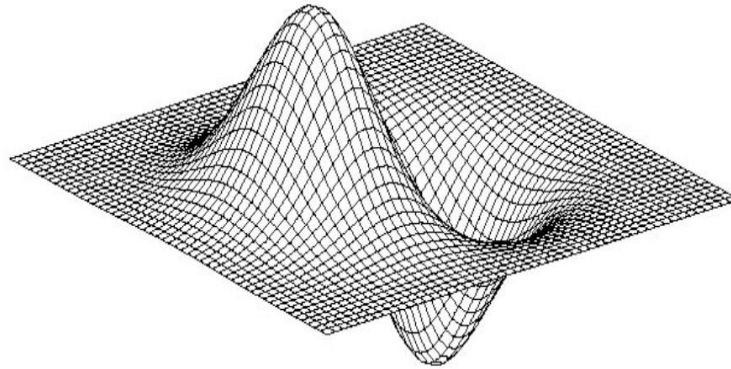
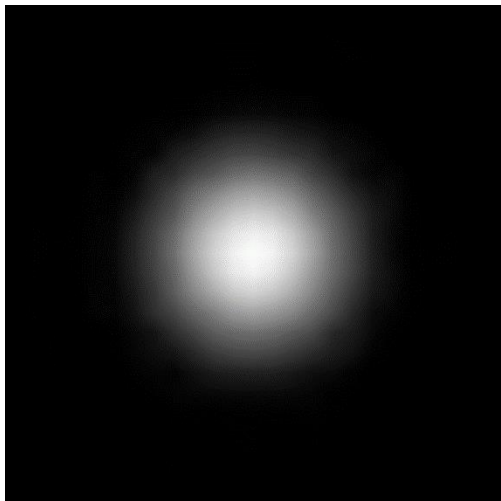
$$\frac{\partial f}{\partial y} = \begin{bmatrix} 1 \\ -1 \end{bmatrix}$$

Gaussian Blur & Its Derivatives

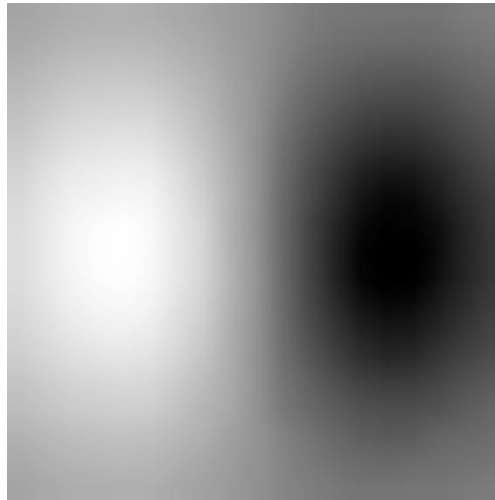
- Blur image & derivative = Convolve with derivative of blur



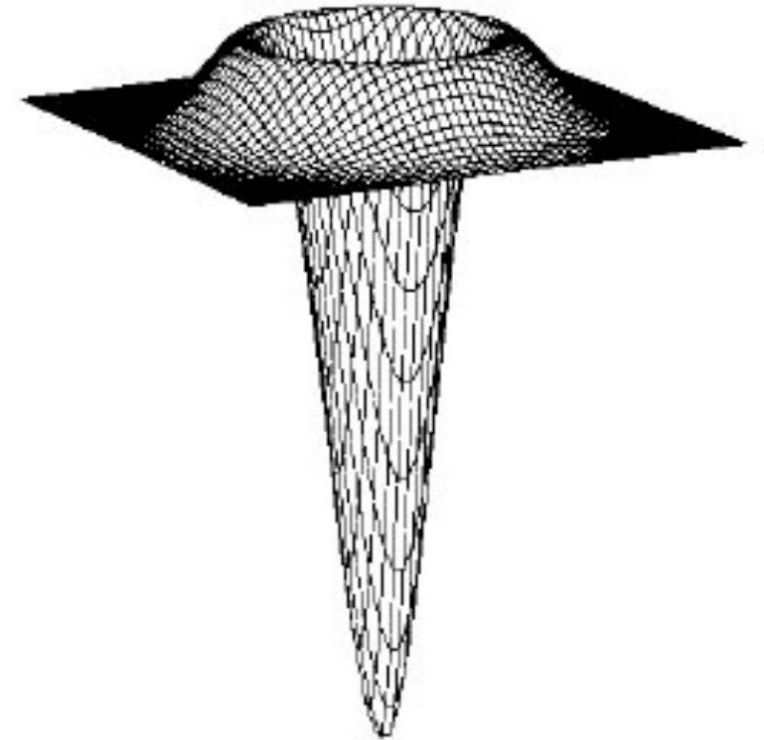
Gaussian



Derivative of Gaussian



Laplacian of Gaussian



Repeated Convolutions: Blur & Derivatives

$$\frac{\partial f}{\partial x} = \frac{1}{8} \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

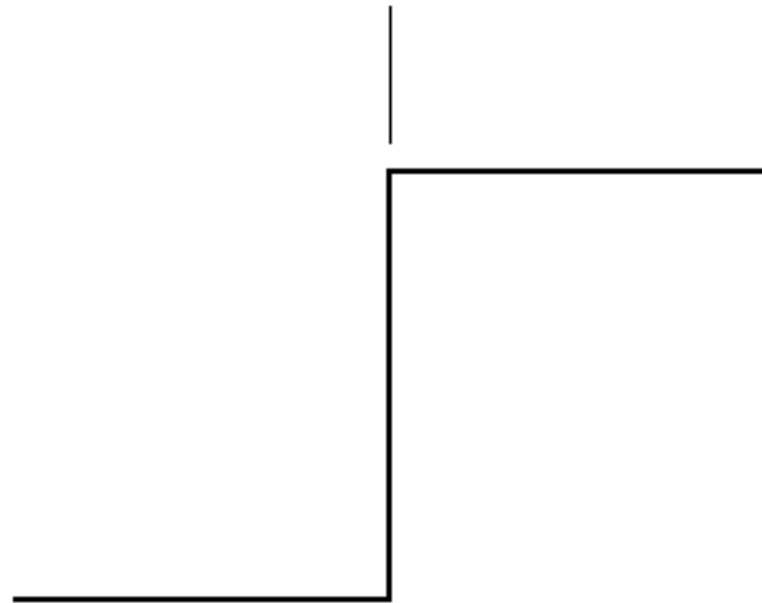
$$\frac{\partial f}{\partial y} = \frac{1}{8} \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

*Home Exercise: Compute **A**, **B**, **C**, and **D***

- $f * \begin{pmatrix} 1/2 & 1/2 \end{pmatrix} * \begin{pmatrix} 1/2 & 1/2 \end{pmatrix} = f * A$
- $f * \begin{pmatrix} 1/2 & 1/2 \end{pmatrix} * \begin{pmatrix} 1 & -1 \end{pmatrix} = f * B$
- $f * \begin{pmatrix} 1/2 & 1/2 \end{pmatrix} * \begin{pmatrix} 1/2 & 1/2 \end{pmatrix}^T = f * C$
- $f * A * B^T = f * D$

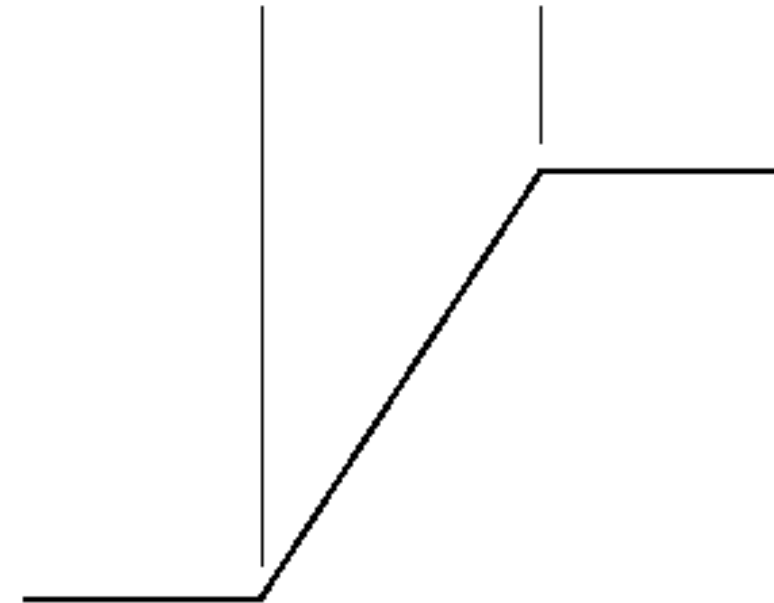
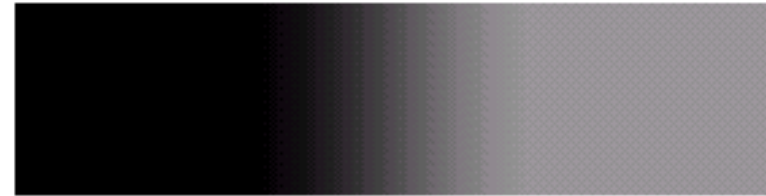
Edges in a picture

Model of an ideal digital edge



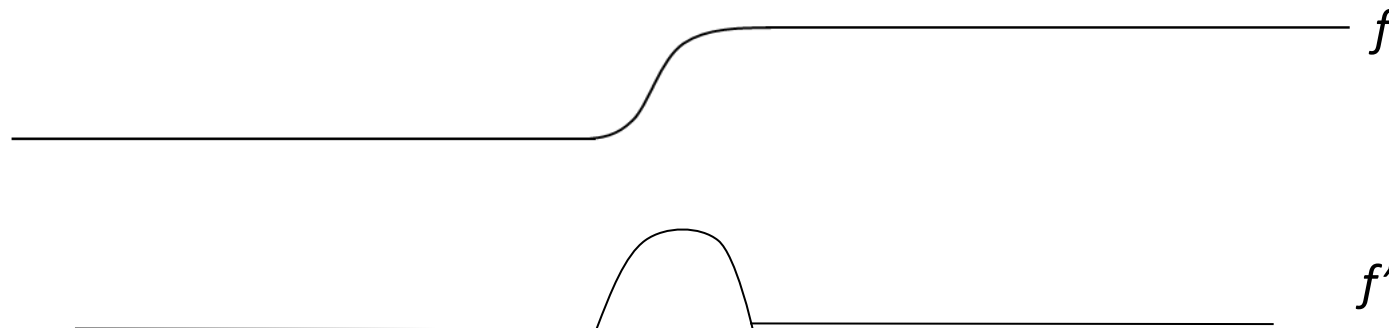
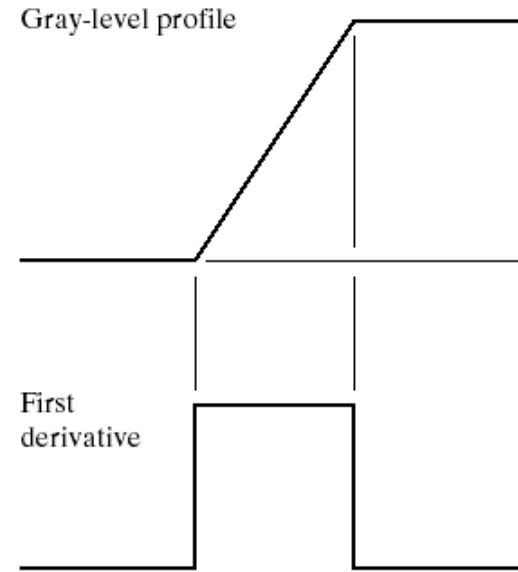
Gray-level profile
of a horizontal line
through the image

Model of a ramp digital edge



Gray-level profile
of a horizontal line
through the image

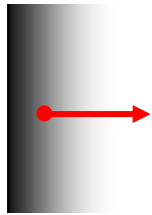
Derivative Response to Edges



The Gradient

Gradient: The vector of x & y derivatives $\nabla f = \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right)$
The direction of most rapid change in intensity

$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

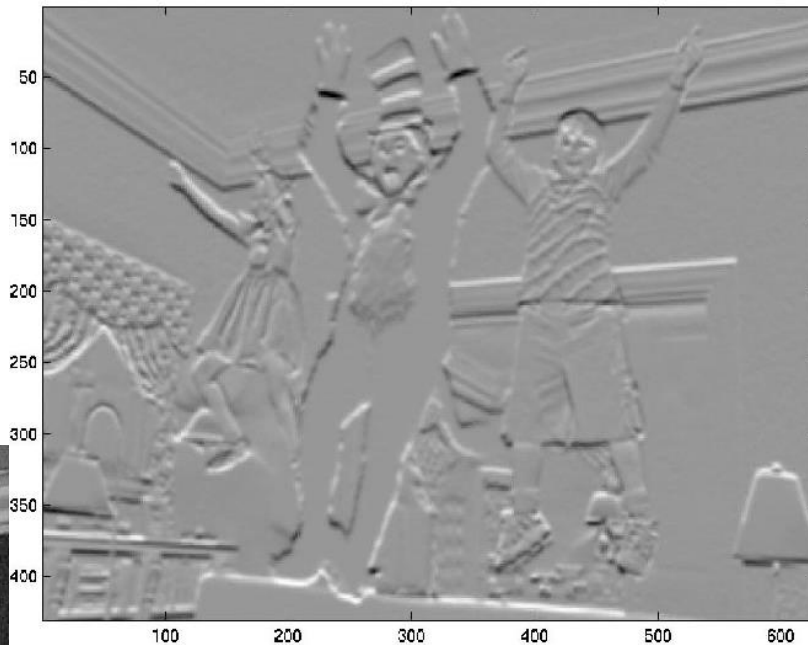
Edge (Gradient) Magnitude

$$|\nabla f| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

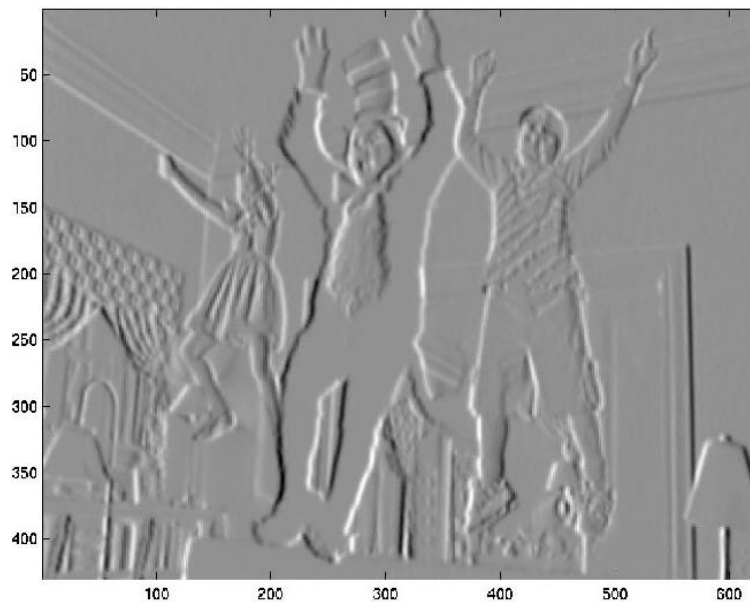
Edge (Gradient) Direction

$$\alpha = \tan^{-1} \left(\left(\frac{\partial f}{\partial y} \right) / \left(\frac{\partial f}{\partial x} \right) \right)$$

$$I_y = \frac{\partial I}{\partial y}$$



Zero is gray



$$I_x = \frac{\partial I}{\partial x}$$

$$|\nabla f| = \sqrt{(I_x)^2 + (I_y)^2}$$



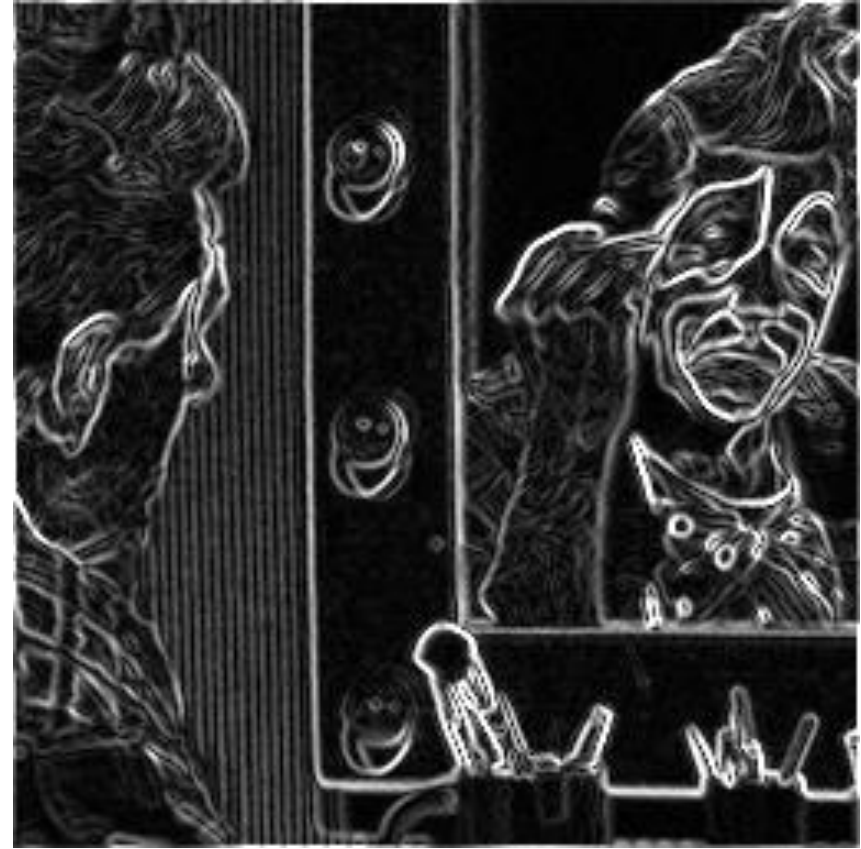
Zero is black

Edge Detection using Gradient

Original



| Gradient |



Second Derivatives

1st x Derivative, convolution with

$$\begin{pmatrix} 1 & -1 \end{pmatrix}$$

2nd x Derivative, convolve again with

$$\begin{pmatrix} 1 & -1 \end{pmatrix}$$

giving

$$\begin{pmatrix} 1 & -2 & 1 \end{pmatrix}$$

1st y Derivative, convolution with

$$\begin{pmatrix} 1 \\ -1 \end{pmatrix}$$

2nd y Derivative, convolve again with

$$\begin{pmatrix} 1 \\ -1 \end{pmatrix}$$

giving

$$\begin{pmatrix} 1 \\ -2 \\ 1 \end{pmatrix}$$

Convolution $\frac{\partial}{\partial^2 x}$: $\begin{pmatrix} 1 & -2 & 1 \end{pmatrix}$

Convolution $\frac{\partial}{\partial^2 y}$: $\begin{pmatrix} 1 \\ -2 \\ 1 \end{pmatrix}$

Laplacian

Equation:

$$\nabla^2 f = \frac{\partial}{\partial x^2} f + \frac{\partial}{\partial y^2} f$$

Convolution:

$$(1 \quad -2 \quad 1) + \begin{pmatrix} 1 \\ -2 \\ 1 \end{pmatrix} = \begin{pmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{pmatrix}$$

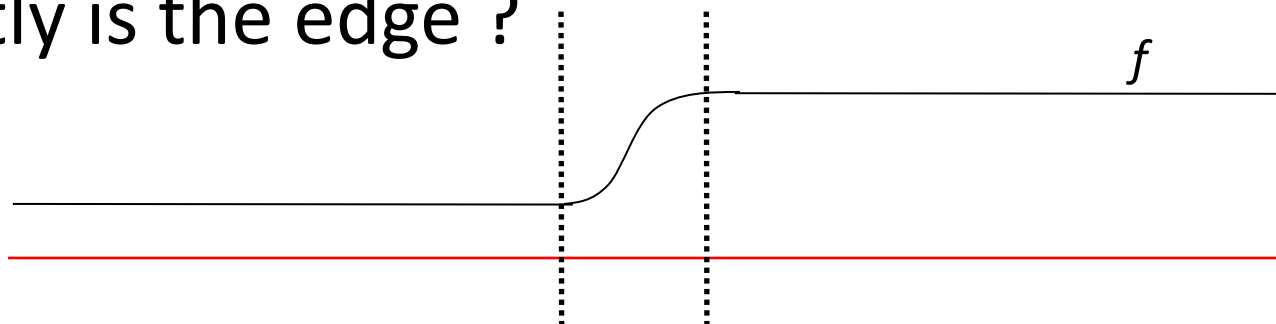
Difference between
a pixel and its
neighborhood

Alternative Laplacian:

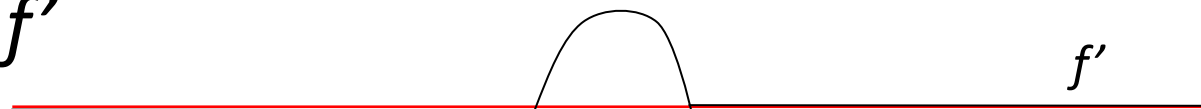
$$\begin{pmatrix} 1 & 1 & 1 \\ 1 & -8 & 1 \\ 1 & 1 & 1 \end{pmatrix}$$

Edge Localization with Zero Crossing

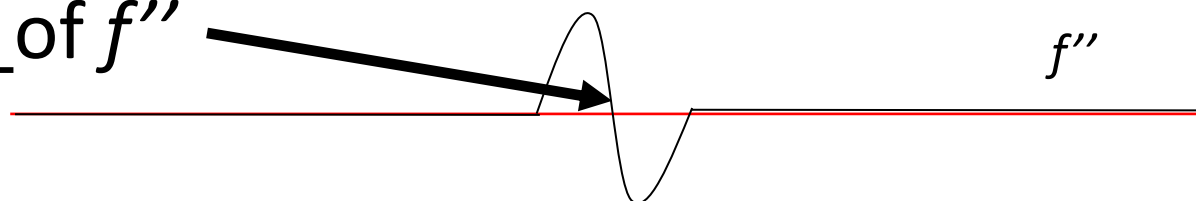
Where exactly is the edge ?



Maximum of f'



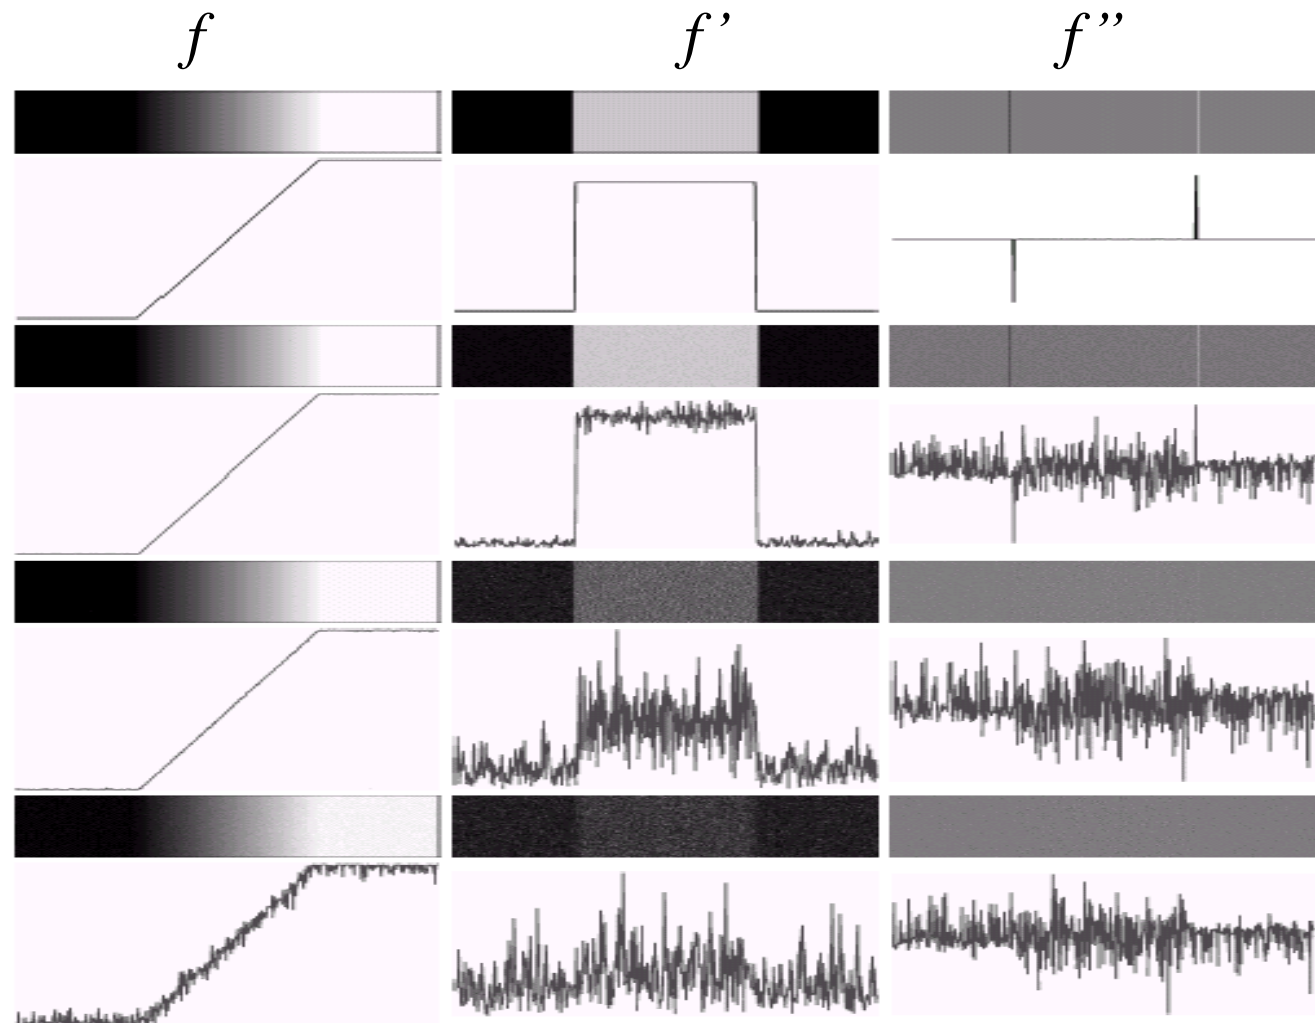
Zero crossing of f''



Problem: f'' is very noisy \Rightarrow Smooth f first !

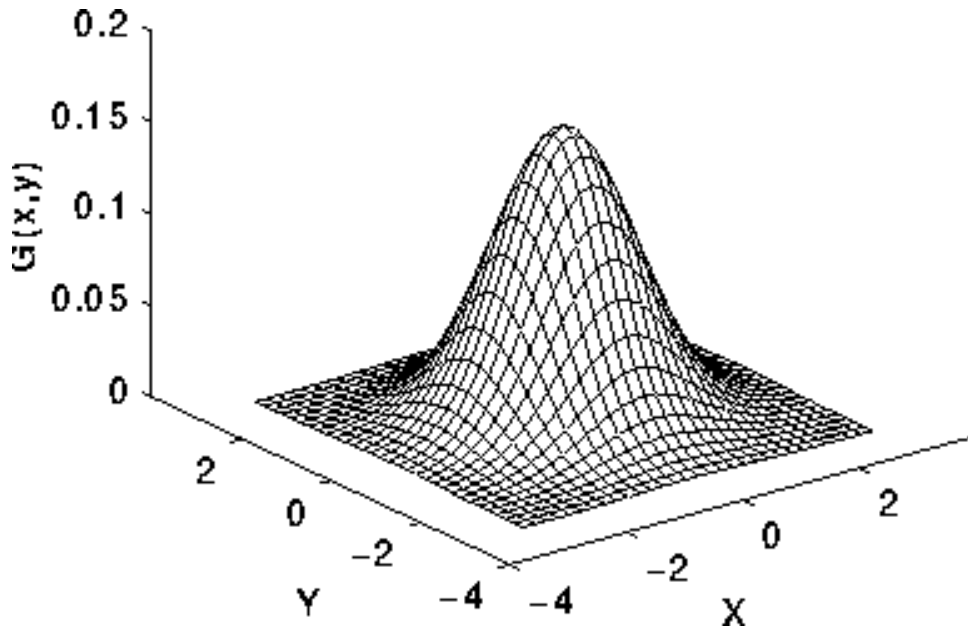
Effect of Noise on Derivatives

Derivatives amplify noise



Smoothing with a 2D Gaussian

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



		1	1		
		1	2	1	
	1	3	3	1	
1	4	6	4	1	

(Gaussian is approximated by binomial coefficients. **Why?**)

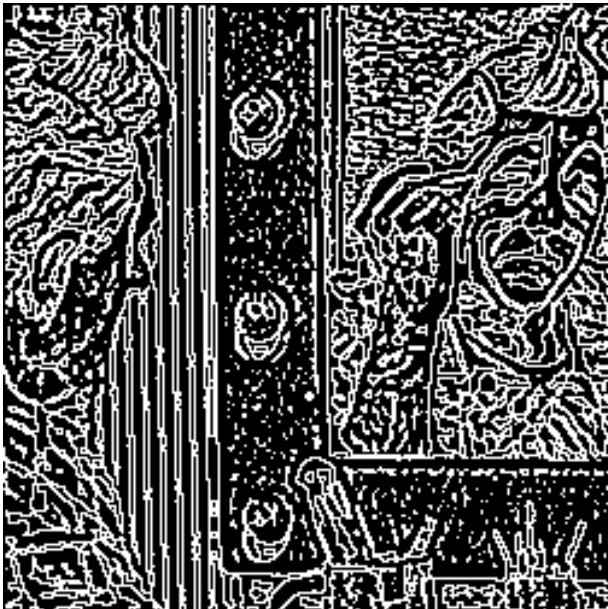


Zero Crossing: Smoothing Effect

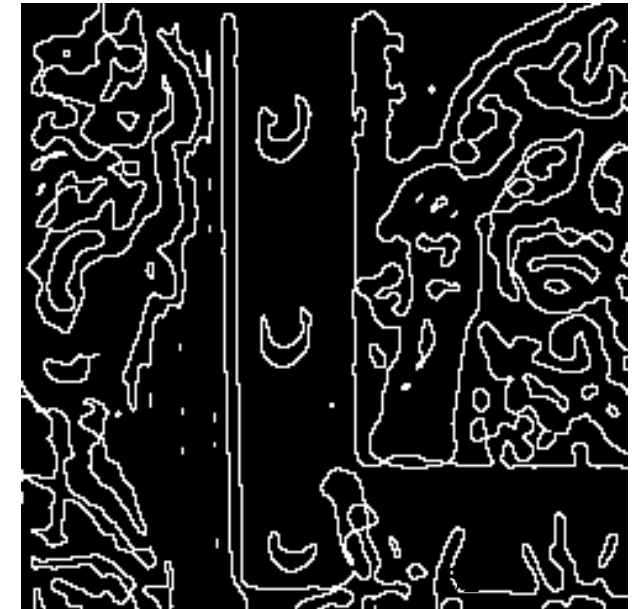
Why are all lines closed curves?

Level Curves – All pixels with same value
Like elevation curves in topological map

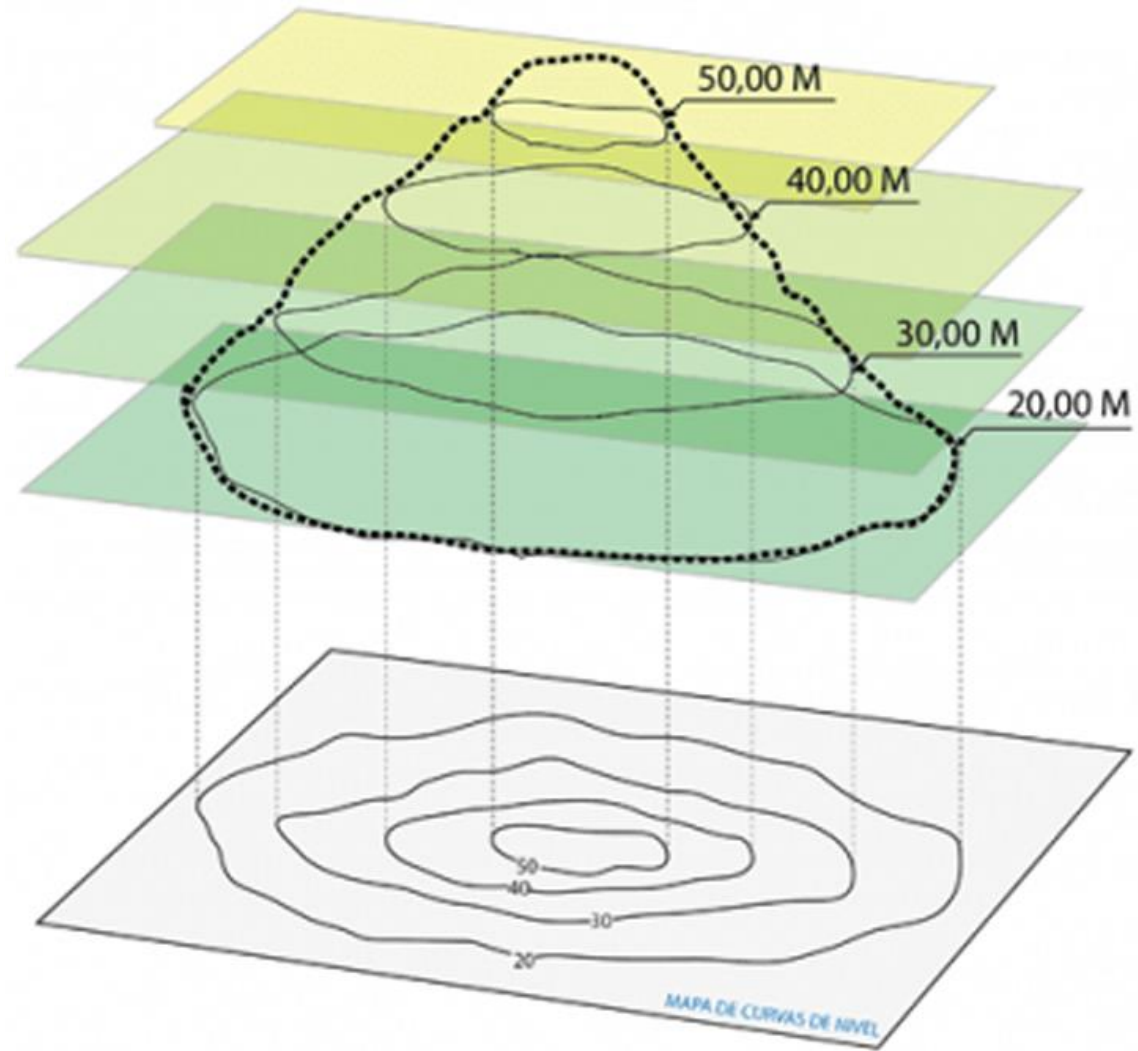
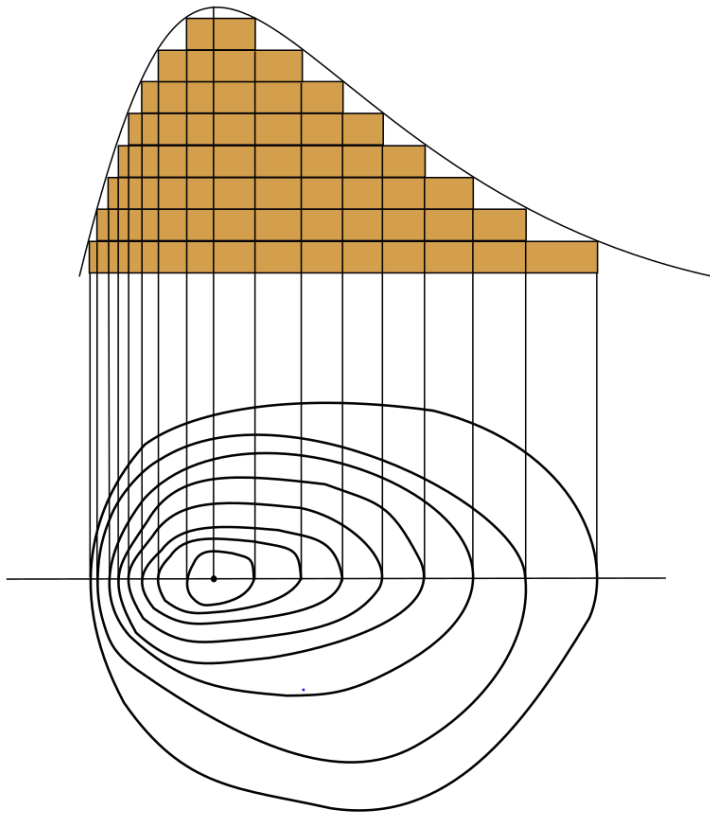
Little Blur



Much Blur



Level Curves - Contour Lines



Canny Edge Detection

- Computing image derivatives f_x, f_y
 - Smoothing with a Gaussian.
 - Using simple derivative kernels $(1, -1)$, $\frac{1}{2}(1, 0, -1)$.
- Computing edge direction: $\tan(\alpha) = f_y / f_x$
- Edge point is the local maxima in edge direction.
 - E.g. zero crossing (to get an edge with width 1)
- Use only edge points with gradient above threshold
- Hysteresis: Edge linking with two thresholds: high and low
 - Low threshold accepted only if neighbor accepted

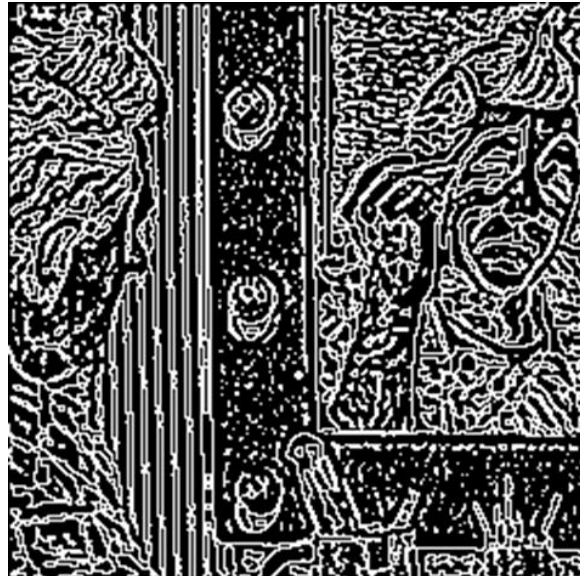


- Blur with different Gaussian widths will give different results

Canny Example



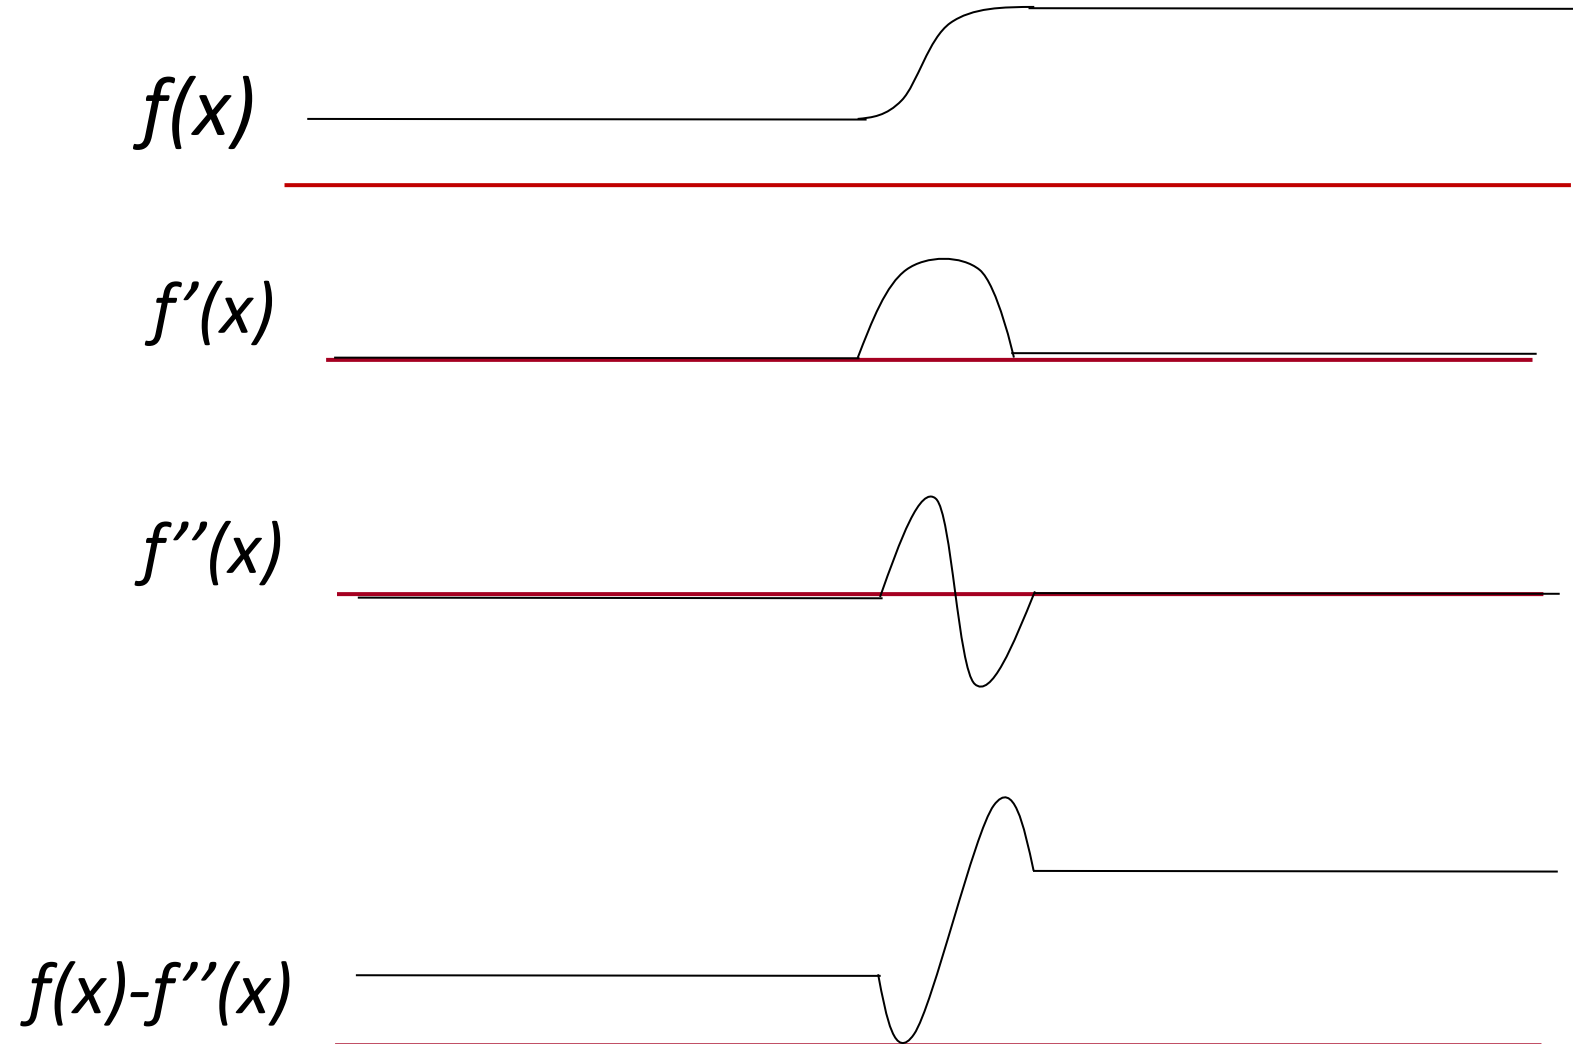
Zero Crossings



Canny



Sharpening by Subtracting the Laplacian



Sharpening by Subtracting the Laplacian

Equation:

$$\nabla^2 f = \frac{\partial^2}{\partial x^2} f + \frac{\partial^2}{\partial y^2} f$$

Convolution:

$$(1 \quad -2 \quad 1) + \begin{pmatrix} 1 \\ -2 \\ 1 \end{pmatrix} = \begin{pmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{pmatrix}$$

Subtracting the Laplacian from the image ($0 < a < 1$): (Check $a = 0.25$)

$$\begin{pmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{pmatrix} - a \begin{pmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{pmatrix} = \begin{pmatrix} 0 & -a & 0 \\ -a & 1 + 4a & -a \\ 0 & -a & 0 \end{pmatrix}$$

Sharpening Example



Convolutional Neural Networks (CNN)

- Historically, people designed convolution kernels for many image processing tasks (“hand crafted features”)
- Recently, CNN’s were introduced to learn convolution kernels from desired input-output data samples.
- A common choice is to keep the kernel size at 3×3 or 5×5 .
- When input has several layers (E.g. RGB = 3 layers), a different kernel is used for each layer, and the results are added together.
- CNN includes different layers, Convolution, Pooling (sampling), activation (e.g. ReLU)