

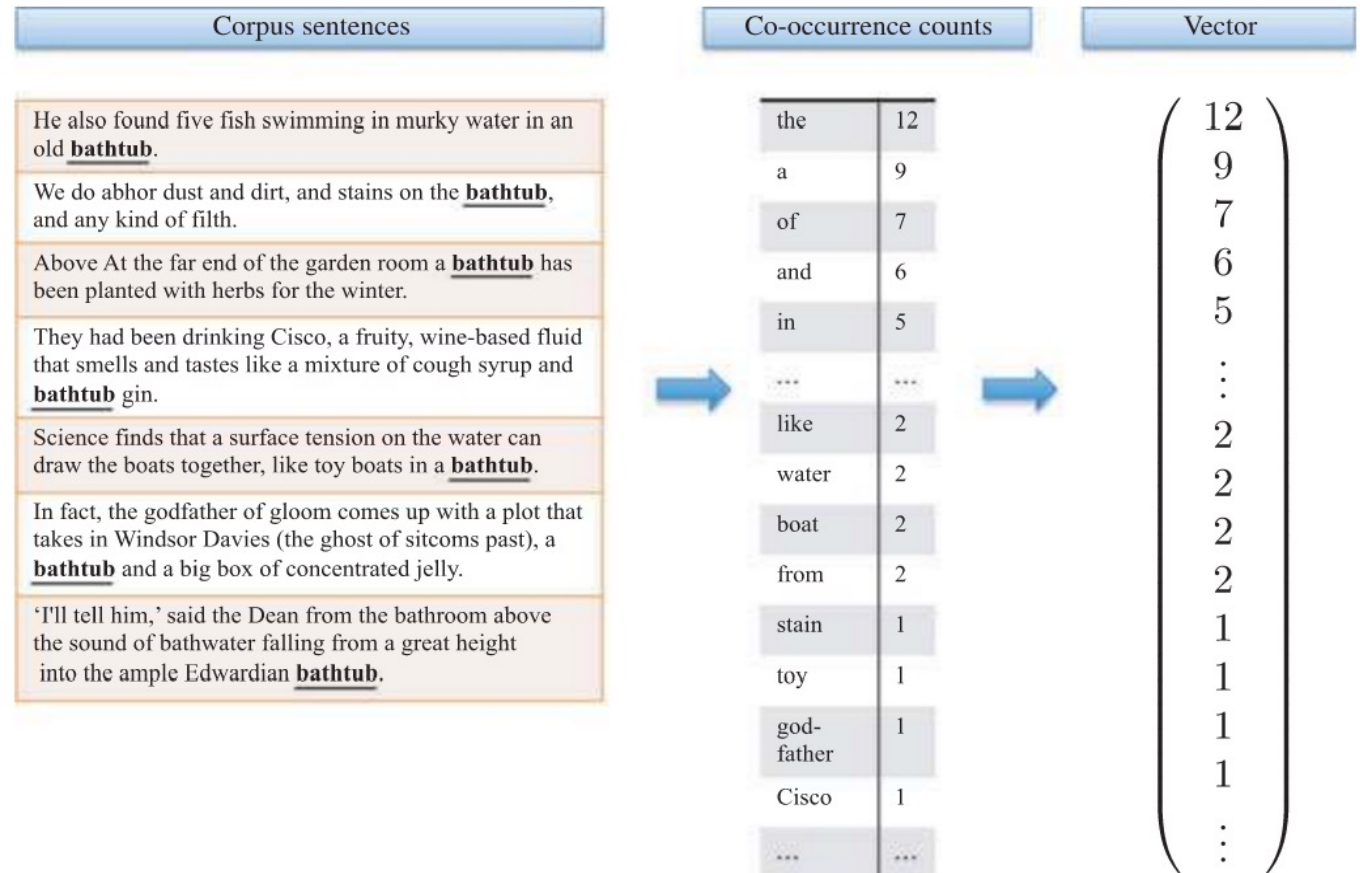
# Lecture 6: Word Embeddings

# The Distributional Hypothesis

- **The distributional hypothesis:** words that are used and occur in the same contexts tend to have similar meanings (Harris, 1954)
- Distributional semantics represents the meaning of words as a distribution over the word's contexts

# Word Embeddings: Count-based Models

- Contexts are defined as neighboring words
  - Usually in a window of  $\pm K$  words
- Dimensions correspond to context wordforms
- Values in entries correspond to counts – the number of times a word and a context word co-occurred




# How Do These Vectors Represent Meaning?

- Similarity can be measured using vector distance metrics
- A popular choice is the “cosine similarity”:

$$\text{similarity}(w, u) = \frac{w \cdot u}{\|w\| \|u\|} = \frac{\sum_{i=1}^n w_i u_i}{\sqrt{\sum_{i=1}^n w_i^2} \sqrt{\sum_{i=1}^n u_i^2}}$$

# How Do These Vectors Represent Meaning?

- Instead of representing words with 1-hot vectors (words are either the same or unrelated), embed them in a space that reflects their similarity patterns
- But taking neighboring wordforms as features is a bit naïve...



# From Word Counts to Dimensionality Reduction

- Distributional semantics makes intuitive sense if we think of the dimensions as representing semantic features
- For example, a *dog* is a mammal, which is terrestrial, a carnivore and often domesticated
  - A *cat* is then more similar to a dog than a goat is, since they share these traits, while goats share only some
- However, neighboring words are considerably less abstract
  - For example, *car* and *automobile* are synonyms; but are represented as distinct dimensions
  - This fails to capture similarity between a word with *car* as a neighbor and a word with *automobile* as a neighbor

# From Word Counts to Dimensionality Reduction

- One way to overcome this would be using dimensionality reduction methods
  - Such as singular value decomposition (Schütze, 1993) or the information bottleneck (Pereira et al., 1993)
- These are strong methods, used (with some variation) today as well
- However, we will devote the rest of the chapter to the more recent, prediction-based models

“Part of Speech Induction from Scratch”, Schütze, 1993

“Distributional clustering of English words”, Pereira, Tishby and Lee, 1993

# Prediction-based Models

- Idea: instead of directly representing the distribution of a word, we can represent words as a vector from which the distribution of a word can be “decoded”
- **Task:** learn a network to predict a neighboring word from a given word
  - Sometimes called “self-supervision”
- An influential suite of methods for prediction-based embeddings is *word2vec* (Mikolov et al., 2013)
- We will review the basic implementation of the skip-gram model



# Skip-gram (Setup)

- Notation:
  - Denote the  $j$ -th wordform in the vocabulary with  $x_j$
  - Denote the vocabulary size (number of distinct wordforms) with  $V$
  - $N$  is a hyperparameter that determines the dimensionality of the embeddings

# Skip-gram (Model)

The probability of predicting the neighbor  $x_j$  given the word  $x_k$

Soft-max over the network's output

$$P(x_j|x_k) = \frac{e^{w'[:,j] \cdot w[k,:]}}{\sum_{j'=1,\dots,V} e^{w'[:,j'] \cdot w[k,:]}}$$


1-hot vector

Distribution over contexts

Input layer

Hidden layer

Output layer



$W$  and  $W'$  are the parameters of the model

# Skip-gram (Training)

- Training is carried out by maximizing

$$\operatorname{argmax}_{W, W'} \log[P(\text{text})] =$$

$$\operatorname{argmax}_{W, W'} \sum_{(x_j, x_k) \in \text{text}} \log[P(x_j | x_k)]$$

where  $x_k$  and  $x_j$  are any pair of words no more than  $K$  tokens apart

# Skip-gram


Each wordform now has two embeddings:

**input embedding** in the input matrix  $W$

- Row  $j$  of the input matrix  $W$  is the  $N$  dimensional embedding for word  $j$  in the vocabulary.

**output embedding**  $v'$ , in output matrix  $W'$


- Column  $j$  of the output matrix  $W'$  is a  $N$  dimensional vector embedding for word  $j$  in the vocabulary.
- The input and output embeddings are often concatenated to yield the word embedding for  $j$



# Word Embeddings as Features

- Word embeddings are often used as features to supervised learning tasks
- For instance, in text classification:

Instead of bag-of-words, one might input the sum of the embeddings of the words in the document



# Word Embeddings as Features

- In this case, we don't really care what the dimensions represent
  - They are just useful feature representations, where the “semantics” of these features are unknown
- In fact, word embeddings are often used as initialization to the first layer of a neural network, and are later updated during training
  - This is sometimes called “fine-tuning”
  - Later in the course...
- Applications cover all aspects of NLP. Essentially any work in NLP that use neural networks (and not only them), will use distributional embeddings to represent the words
  - We will see an example next lesson, and more towards the end of the course